

# How to make the best use of Live Sessions

---

- Please login on time
- Please do a check on your network connection and audio before the class to have a smooth session
- All participants will be on mute, by default. You will be unmuted when requested or as needed
- Please use the “Questions” panel on your webinar tool to interact with the instructor at any point during the class
- Ask and answer questions to make your learning interactive
- Please have the support phone number (US : 1855 818 0063 (toll free), India : +91 90191 17772) and raise tickets from LMS in case of any issues with the tool
- Most often logging off or rejoining will help solve the tool related issues

# COURSE OUTLINE



## Module 02

Introduction to Kubernetes

Kubernetes Architecture

Deploy app to Kubernetes Cluster

Expose App, Scale App And Update App

Managing State with Deployments

Federations, Auditing and Debugging Kubernetes, Security best practices

edureka!



# Kubernetes Architecture

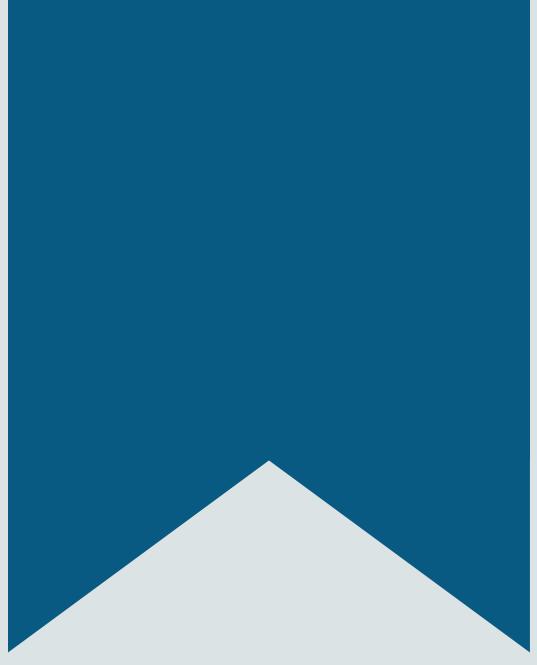
# Objectives

---

After completing this module, you should be able to understand:

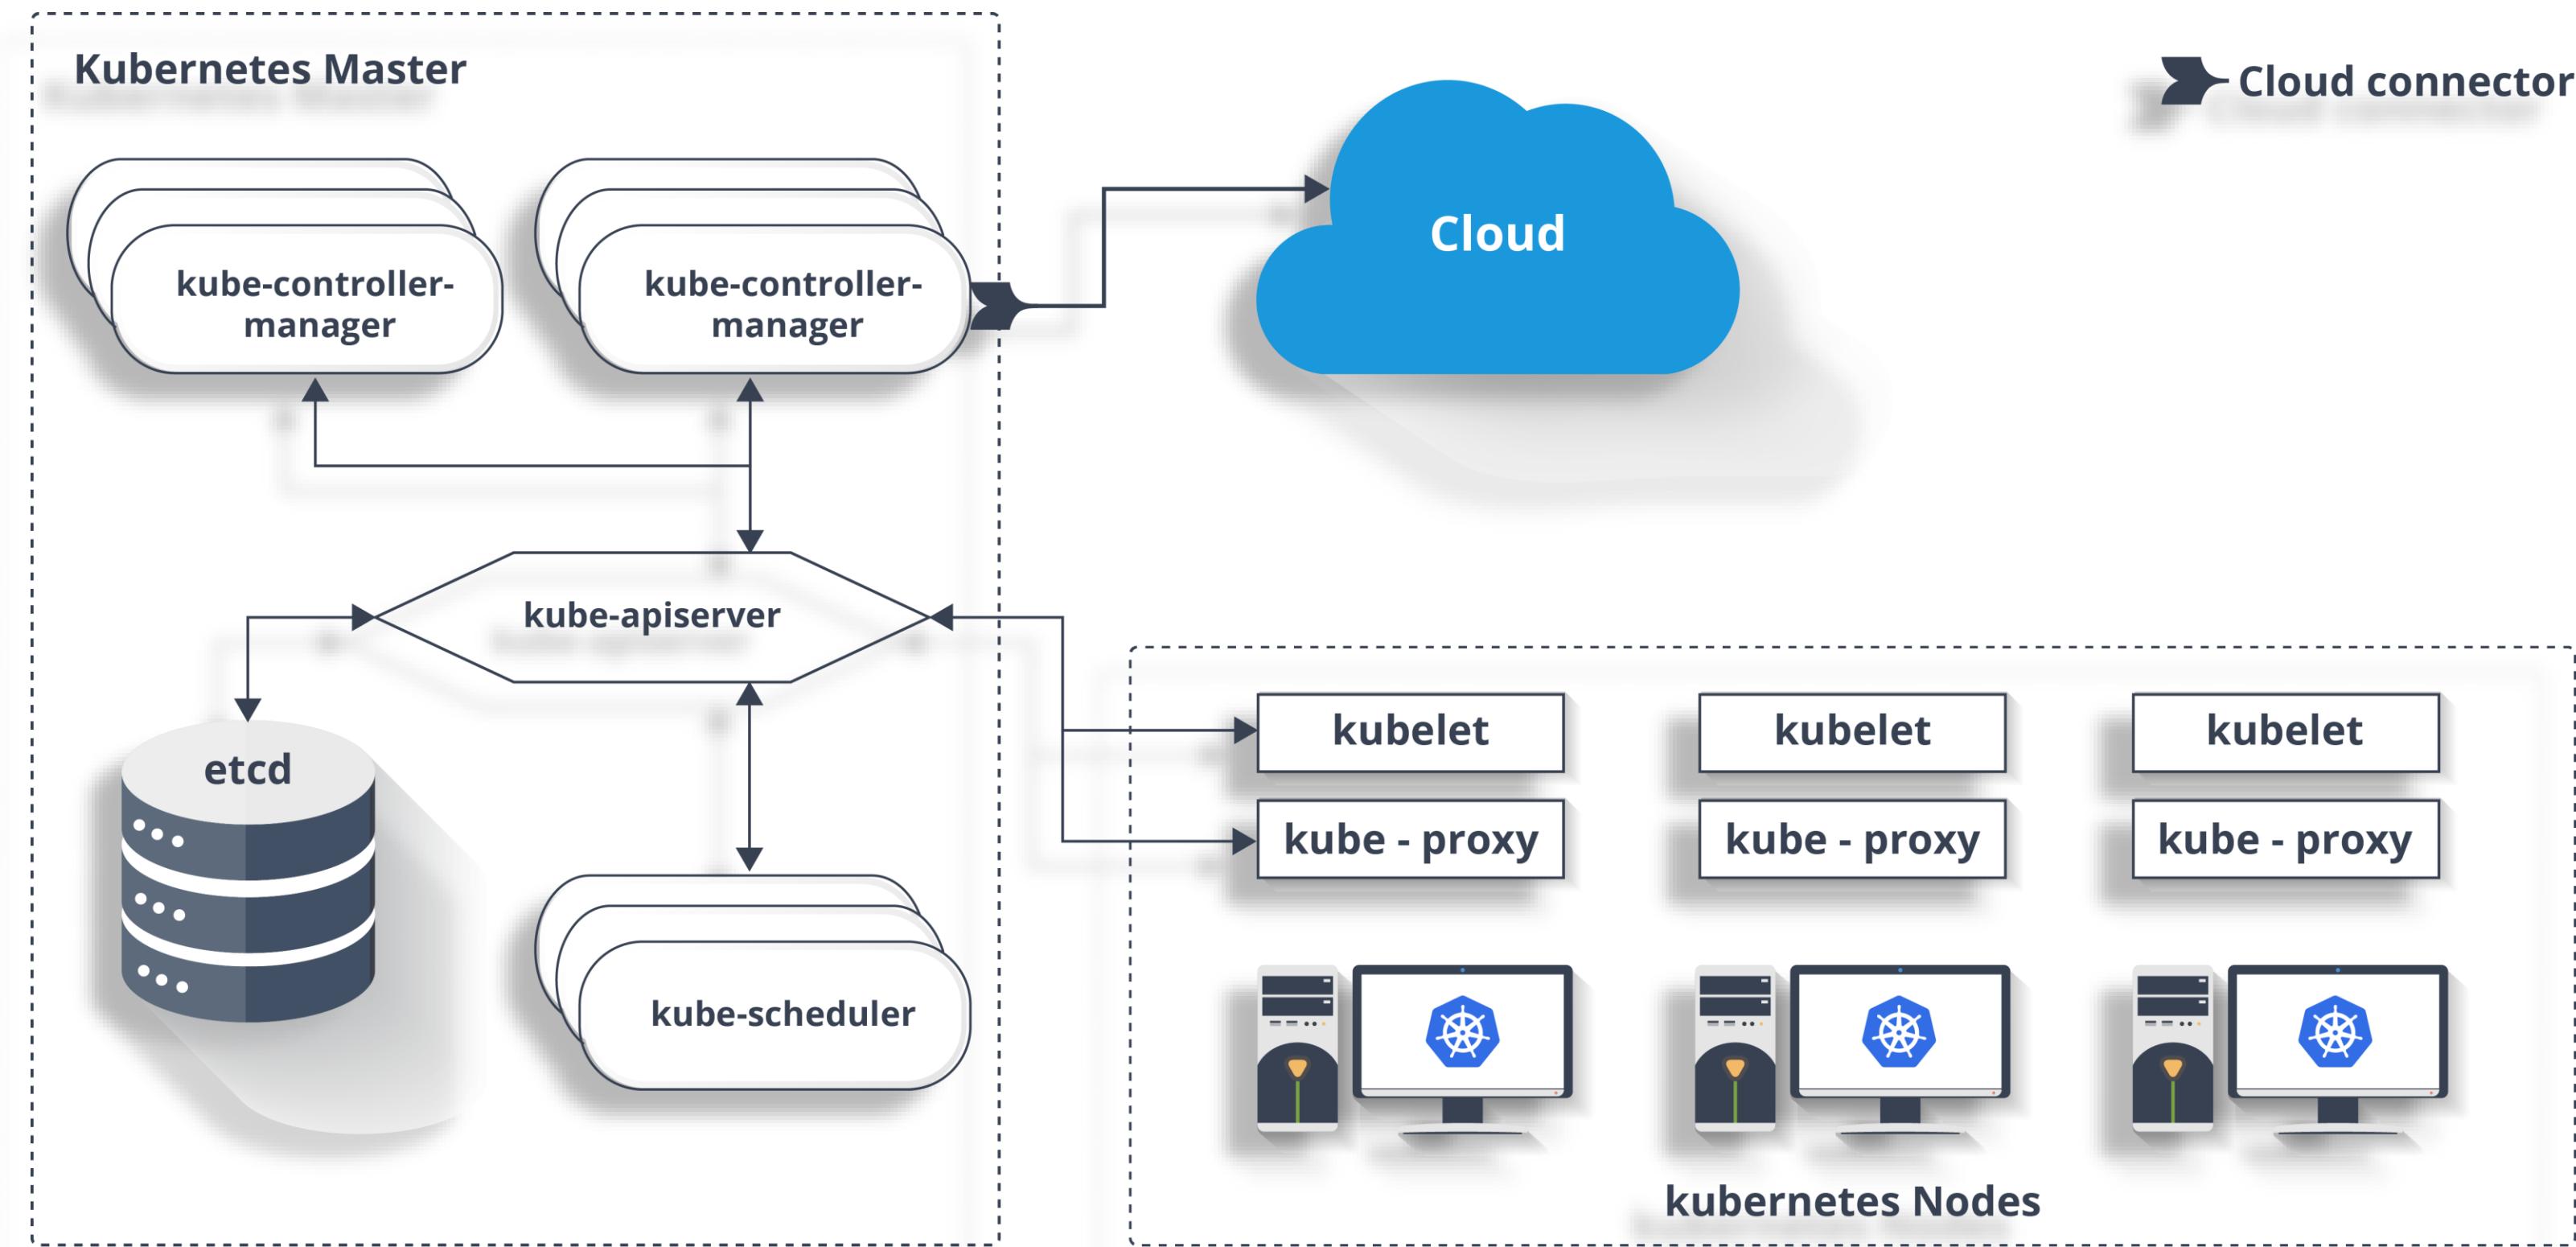
- Master Components of Kubernetes
- Node components of Kubernetes
  - Docker, kubelet, kube-proxy, kubectl
- Addons
  - Cluster DNS
  - Kubernetes Dashboard
  - Container Resource Monitoring
  - Cluster level logging

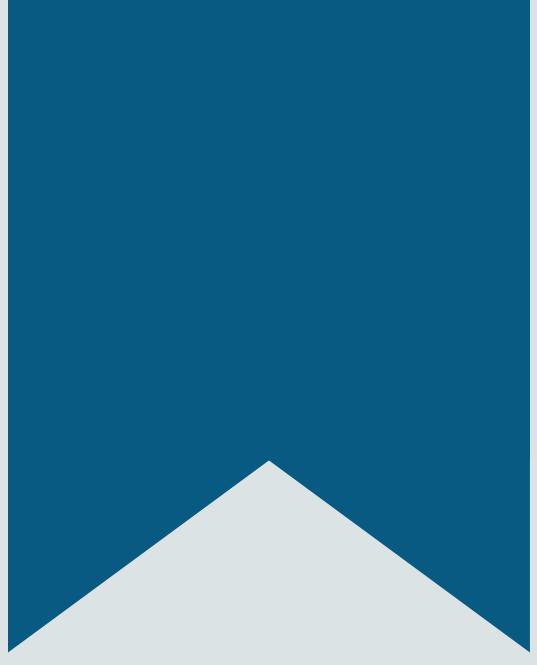




# Kubernetes Architecture

# Architecture Diagram

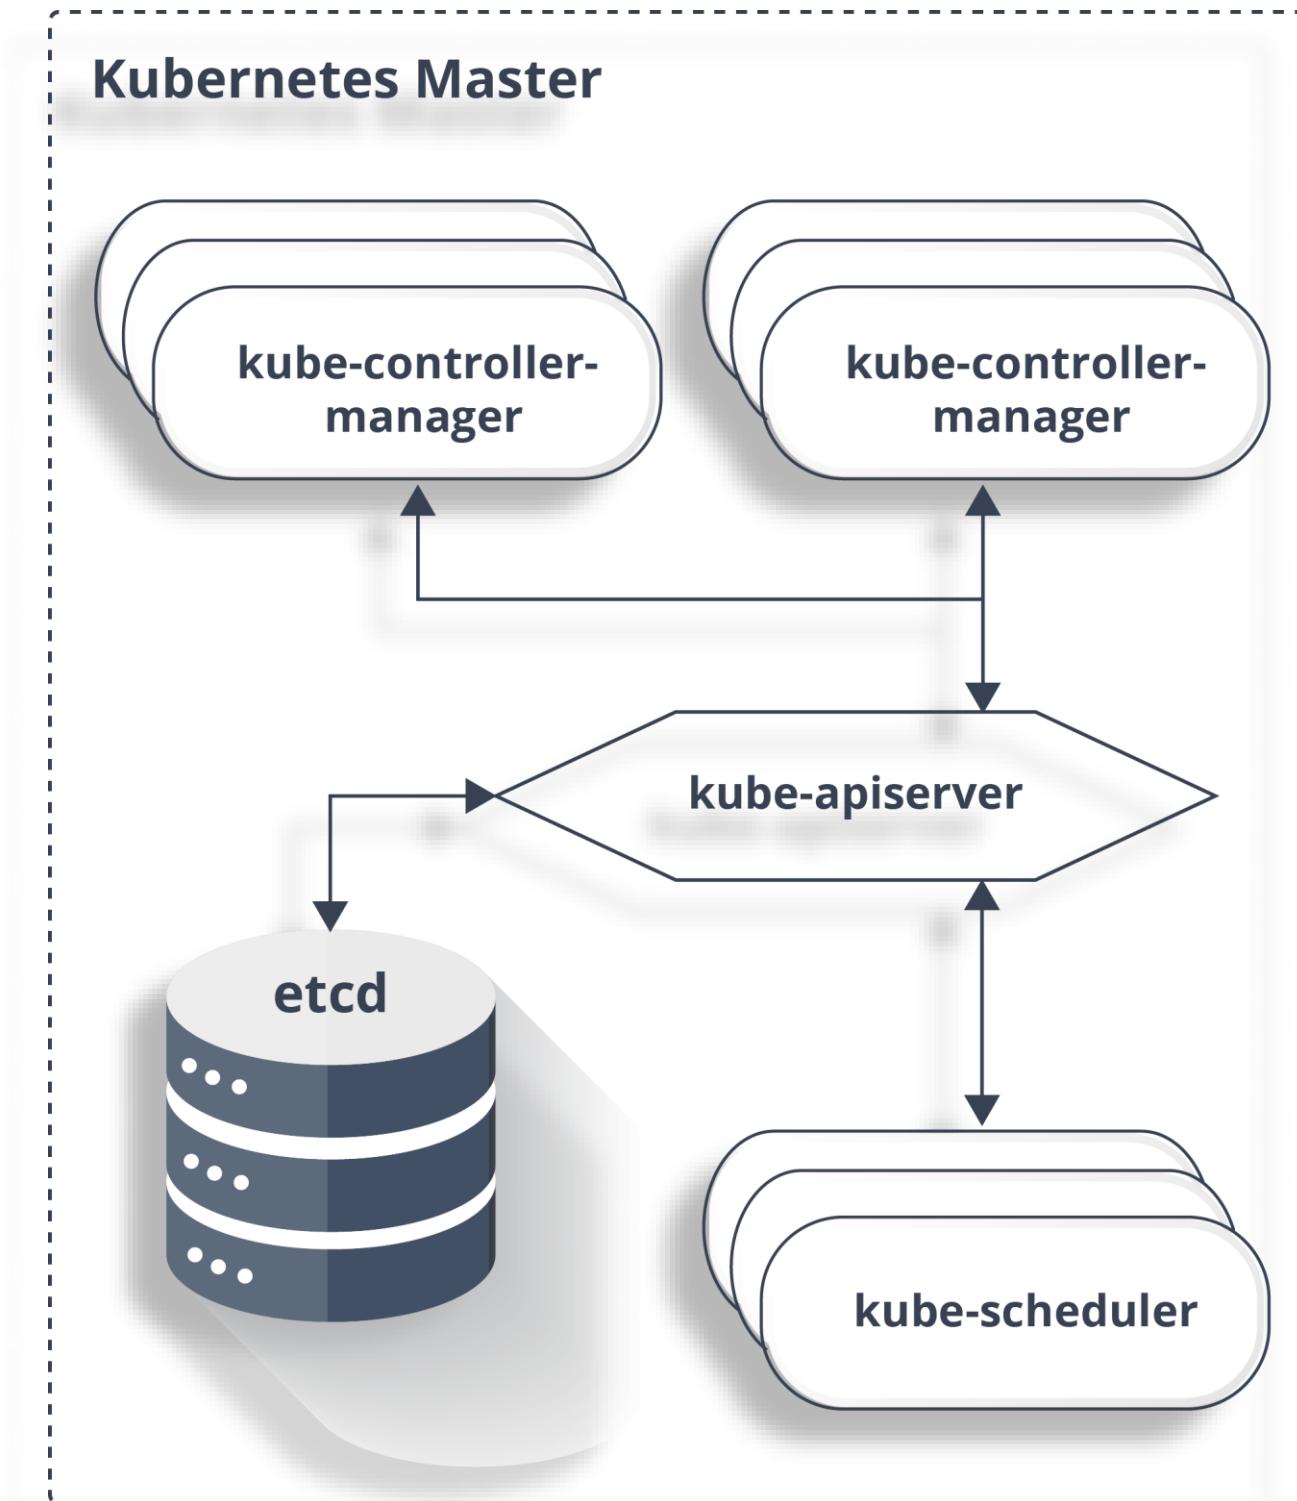
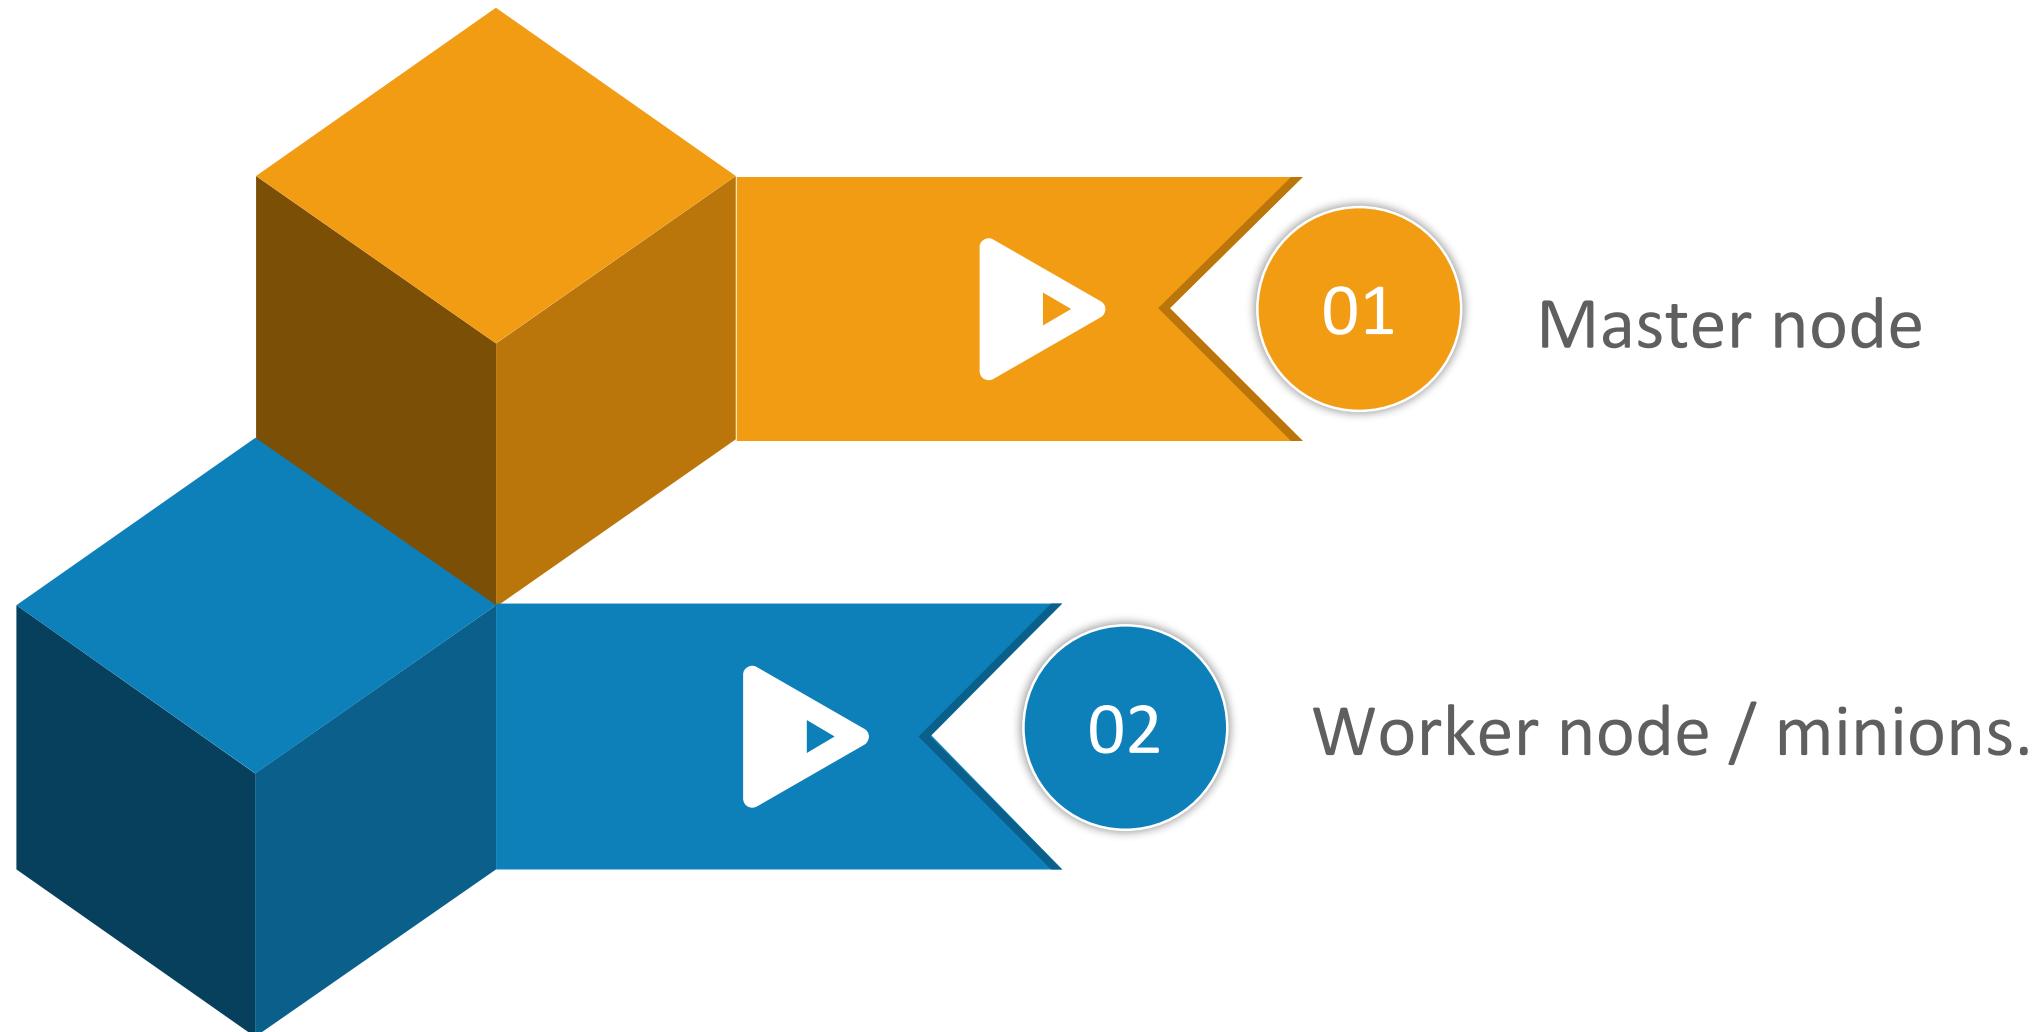




# Master Components of Kubernetes

# Master Components of Kubernetes

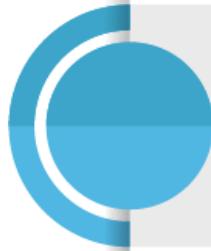
Kubernetes solution is based on two main types of nodes:



# Master Components of Kubernetes



This is the main part of the Kubernetes cluster



Takes all the decision related to the cluster



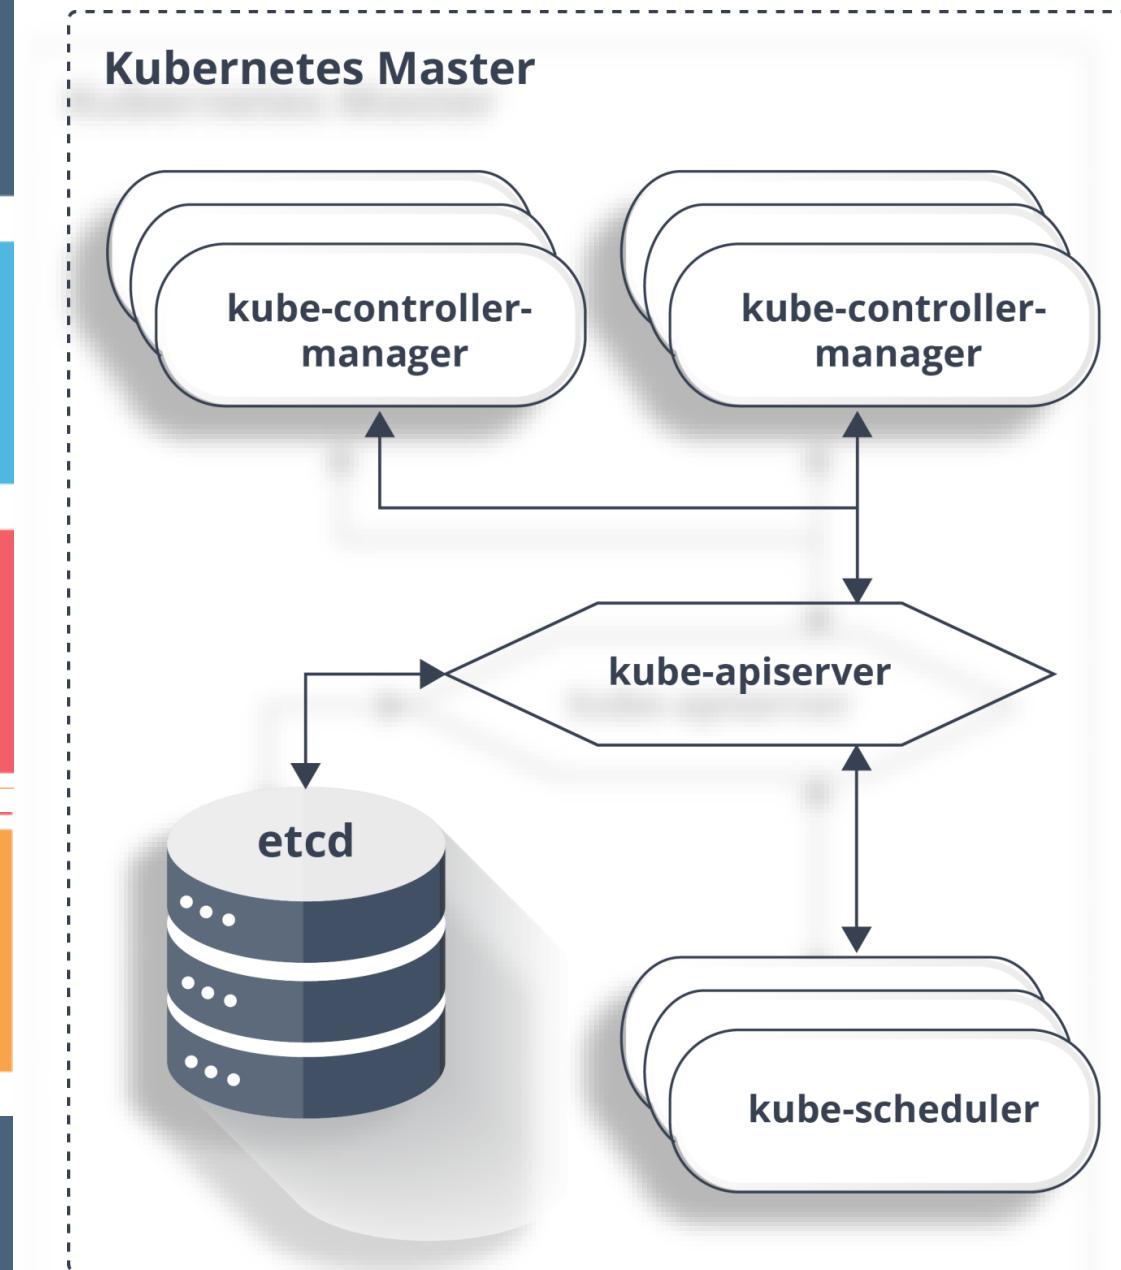
Can be configured on standalone machine



Multiple master are configured in cluster for highly available environment.



As a best practise, don't run user-container on Master node.

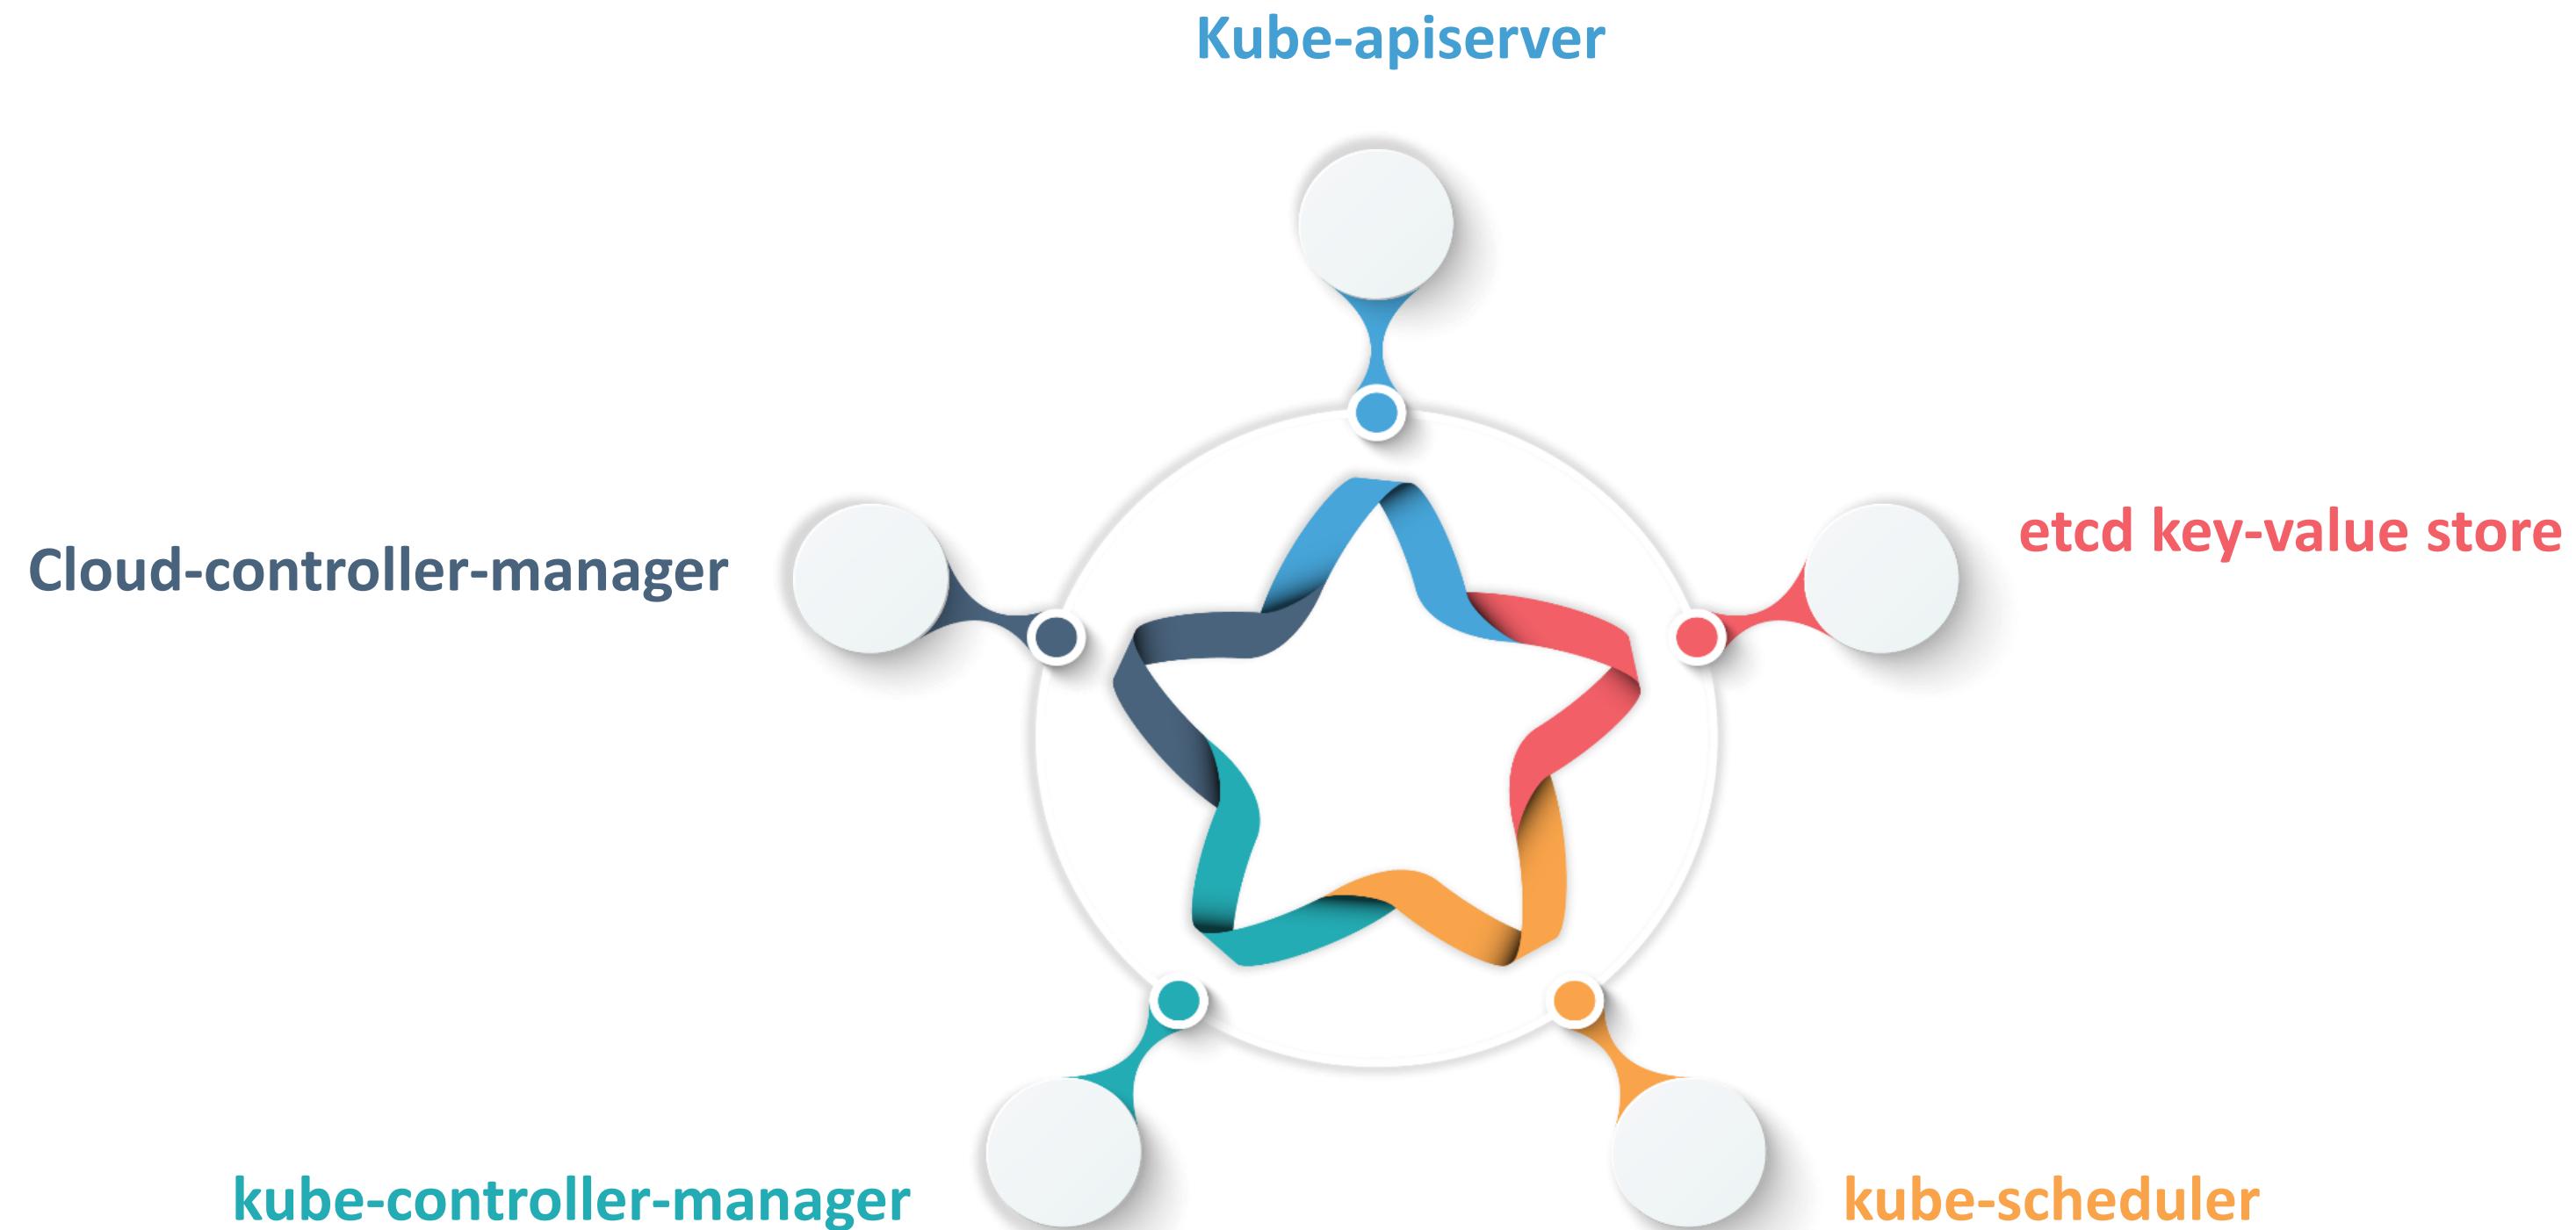


# Master Node Functionality

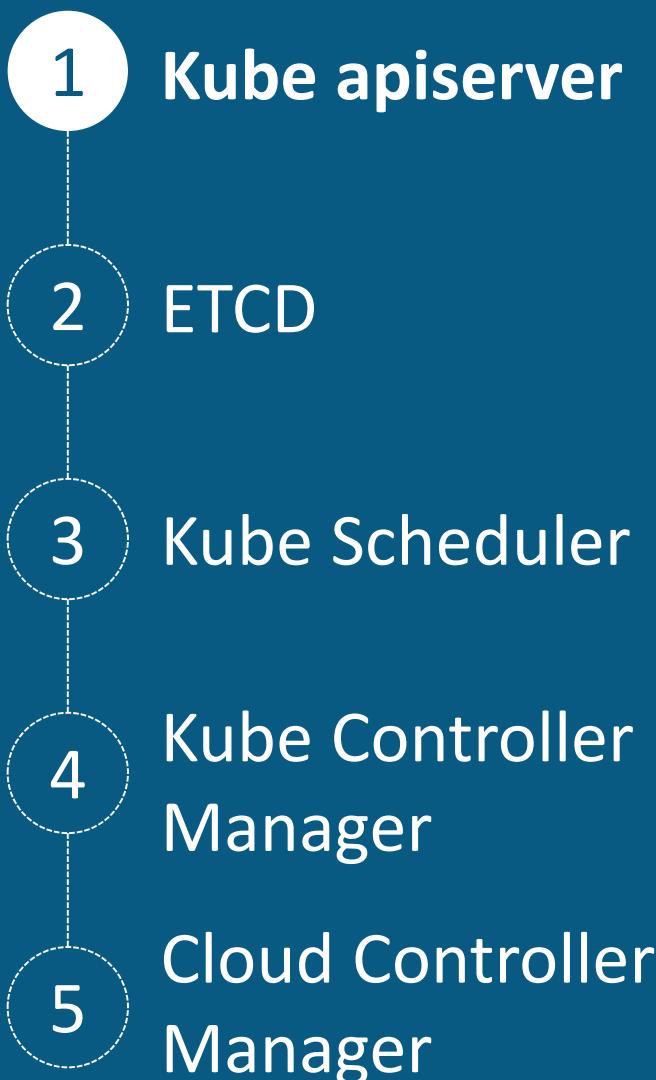


# Master Node Components

---

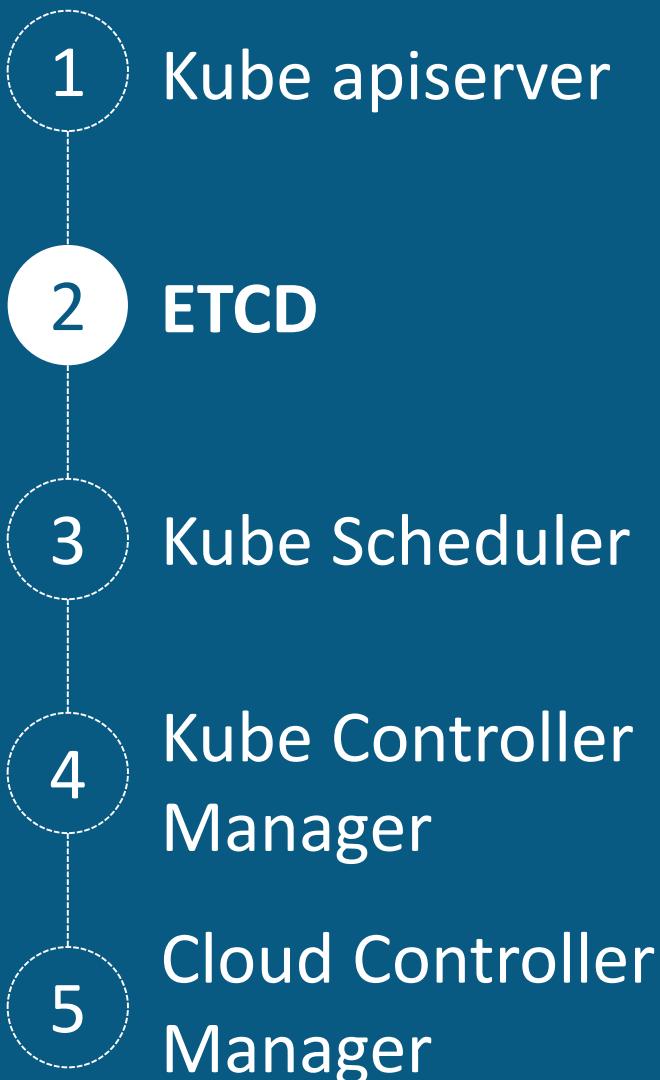


# Master Node Components



- This is the front-end of the master node control plane
- It exposes all the APIs of the kubernetes Master node components
- It is responsible for establishing communication between Kubernetes Node and the Kubernetes Master components
- It follows the scale-out architecture i.e. it can expand horizontally by adding more instances

## Master Node Components



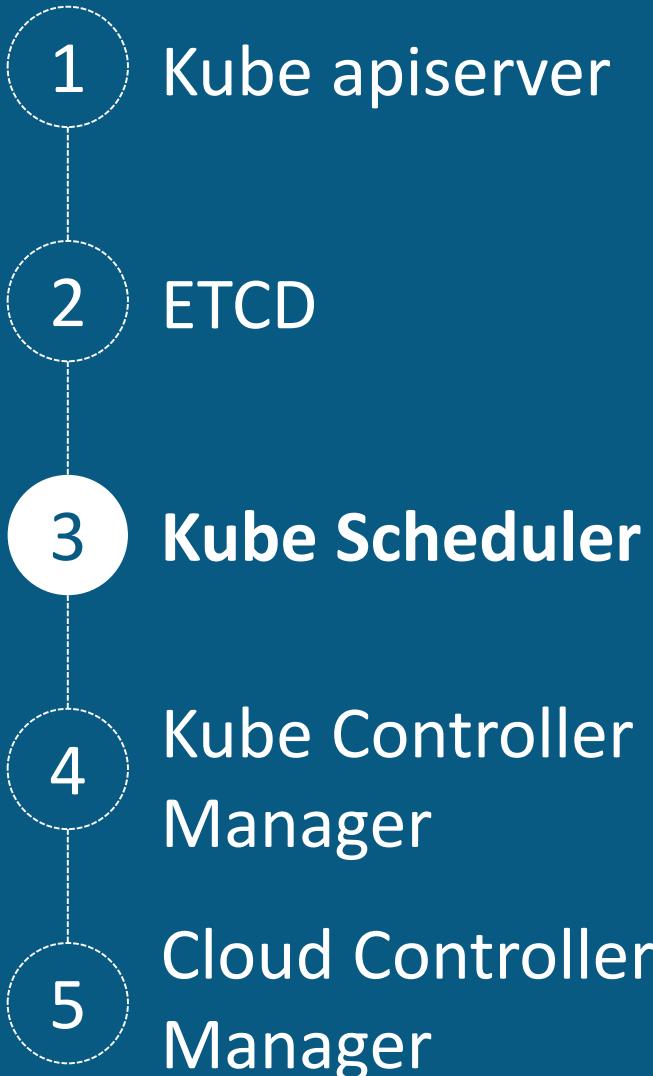
- Kubernetes master node uses key-value store for backing all data of cluster
- “etcd key-value” is distributed datastore which holds all the configuration data of the cluster
- Etcd Key-value is the reference point for all the components in the cluster

## Master Node Components

- 1 Kube apiserver
- 2 ETCD
- 3 Kube Scheduler
- 4 Kube Controller Manager
- 5 Cloud Controller Manager

- Provides reliable datastore through distributed locking mechanism, write barriers and ensure that leader is sending heartbeats periodically
- It is intended to provide a highly available and permanent data storage and retrieval
- It is recommended to have a backup of etcd file

## Master Node Components

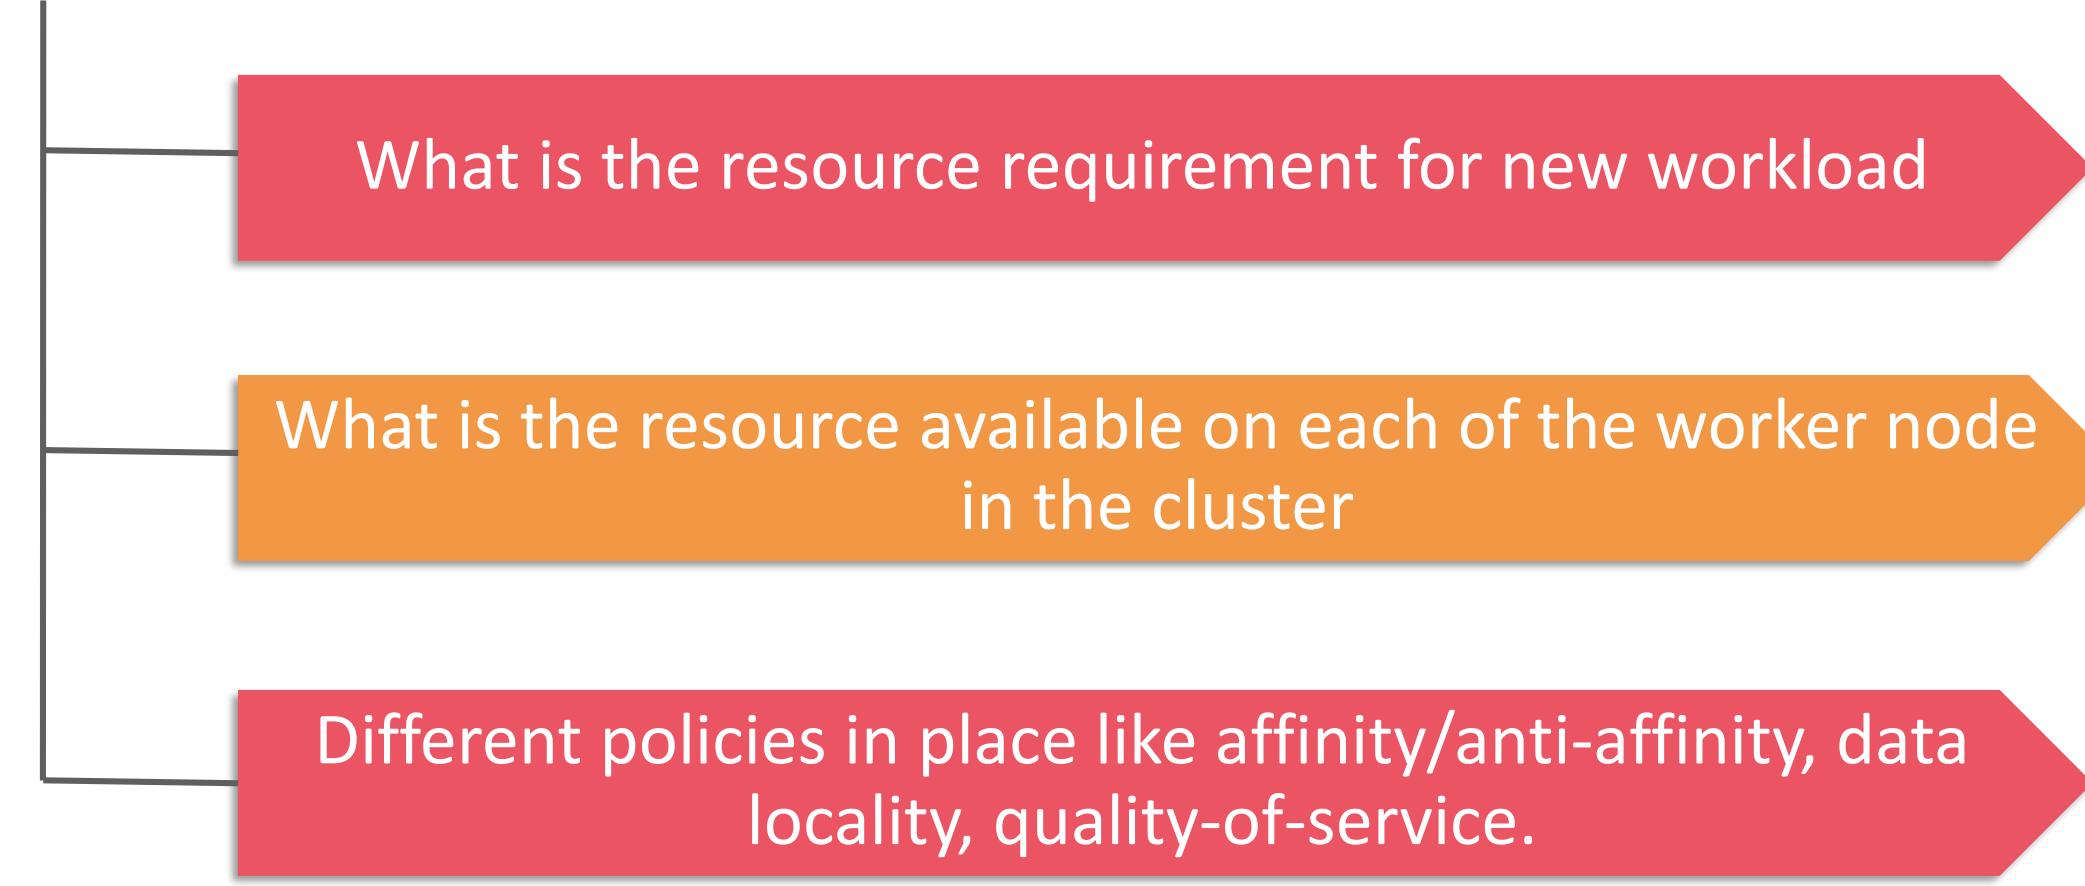


- Responsible for distribution and management of workload on the worker nodes
- It selects the most suitable node to run the unscheduled pod based on resource requirement
- It keeps track of resource utilization and makes sure that workload is not scheduled on nodes which are already full

# Kube-Scheduler

---

For effective resource scheduling, scheduler knows :

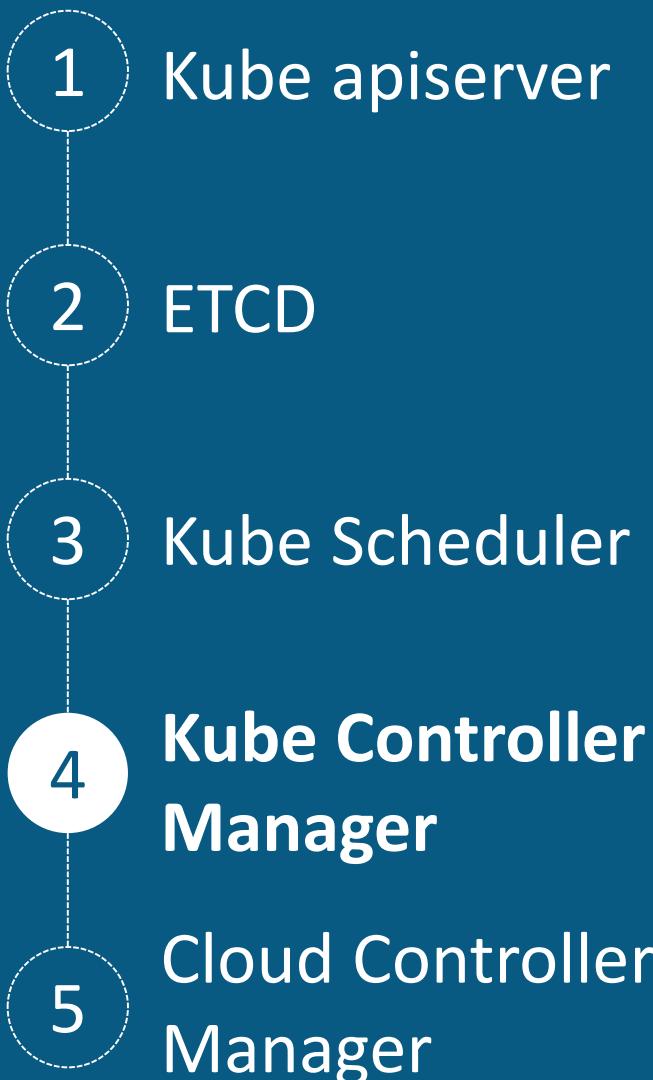


What is the resource requirement for new workload

What is the resource available on each of the worker node  
in the cluster

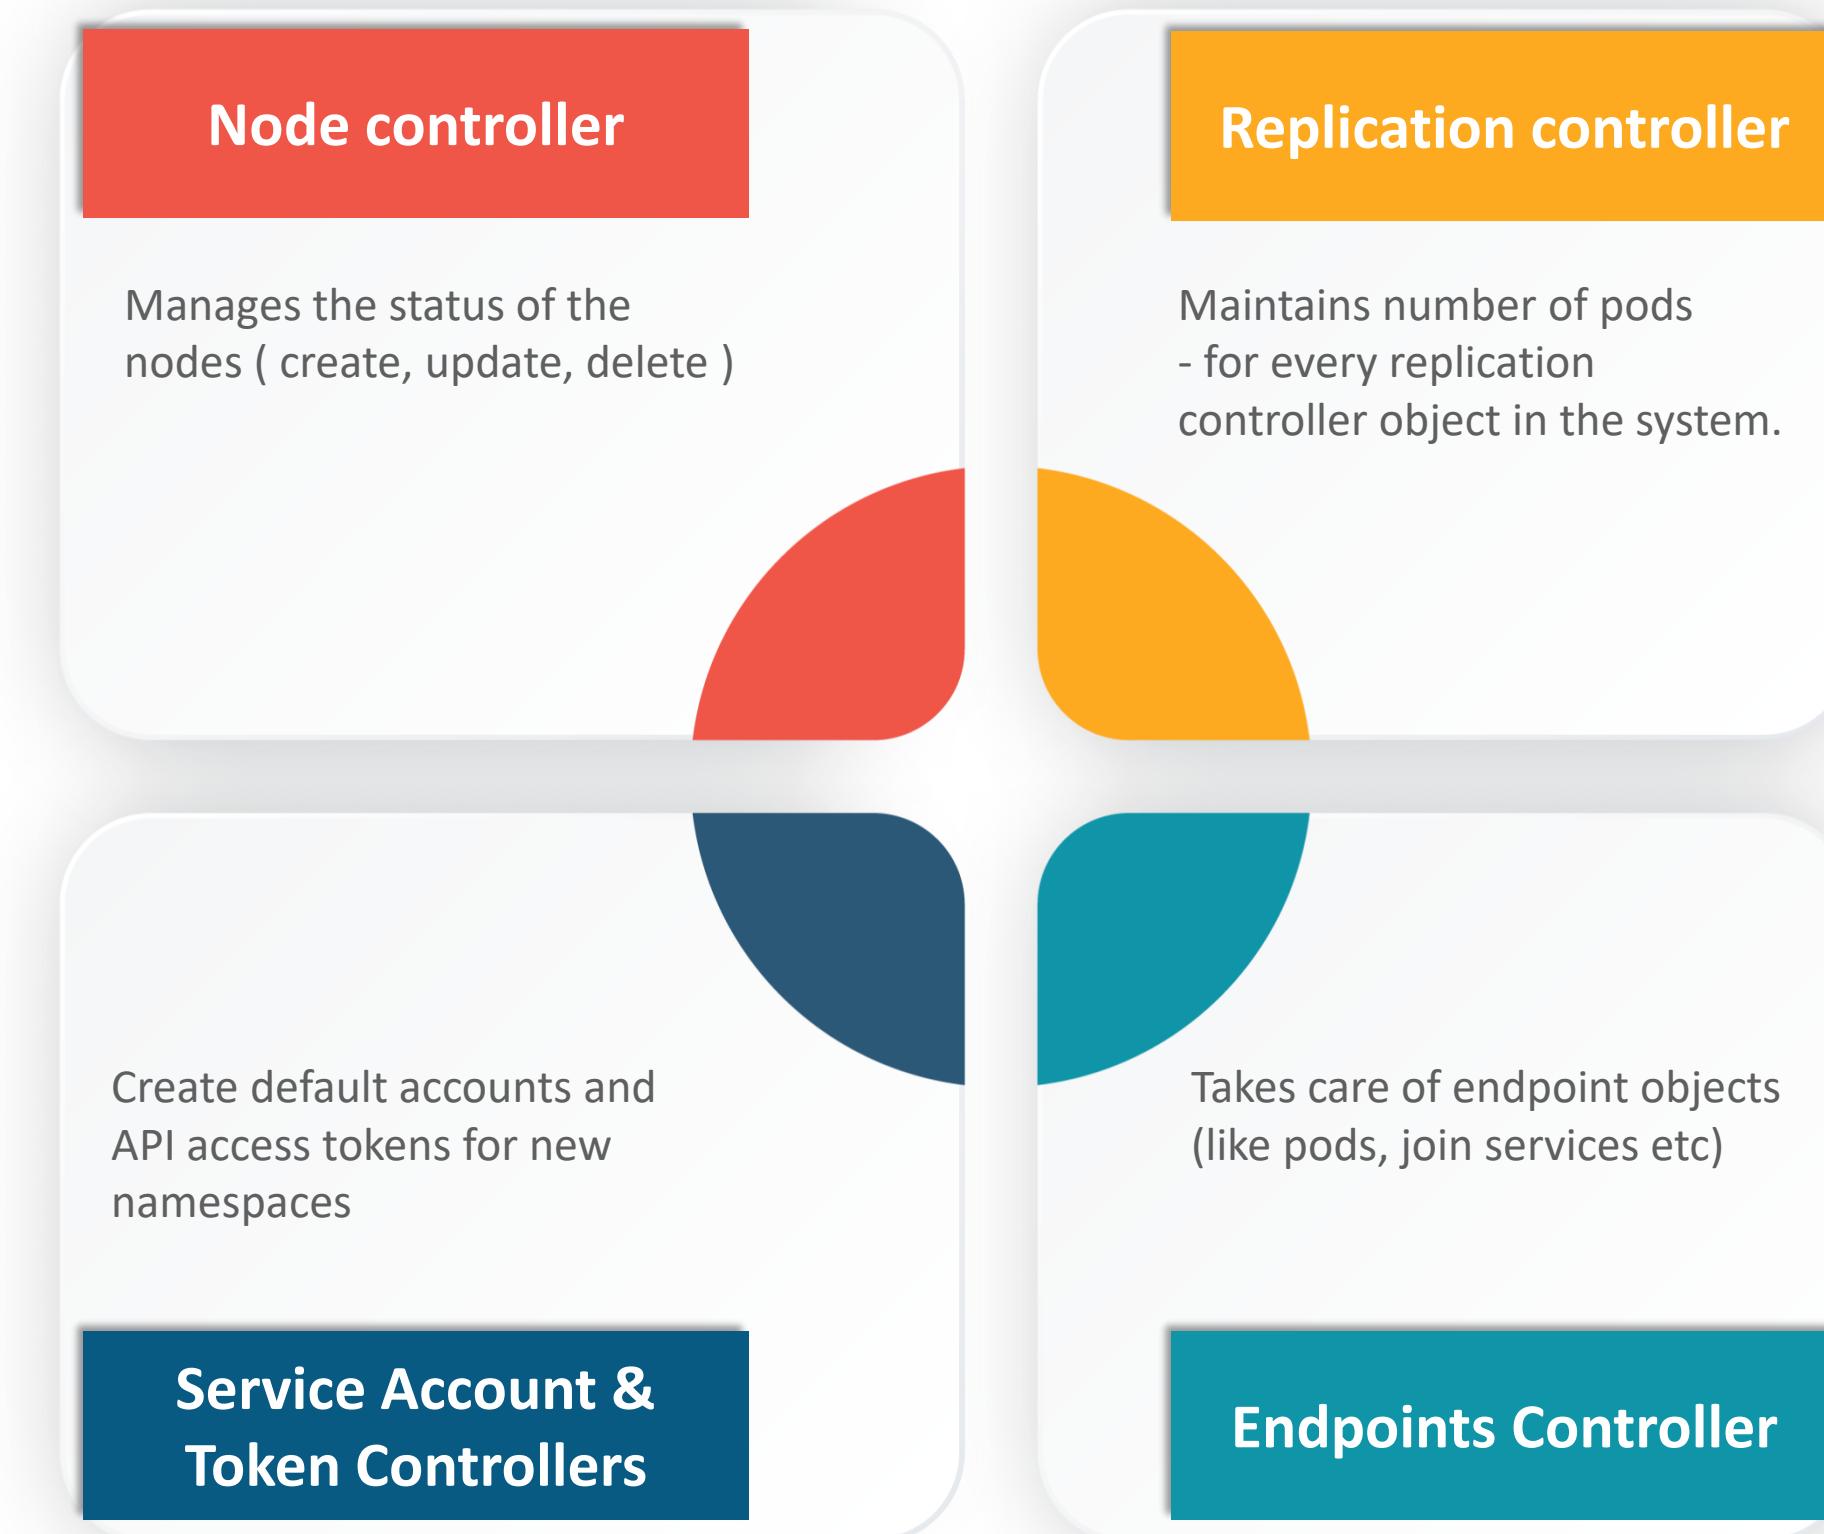
Different policies in place like affinity/anti-affinity, data  
locality, quality-of-service.

## Master Node Components

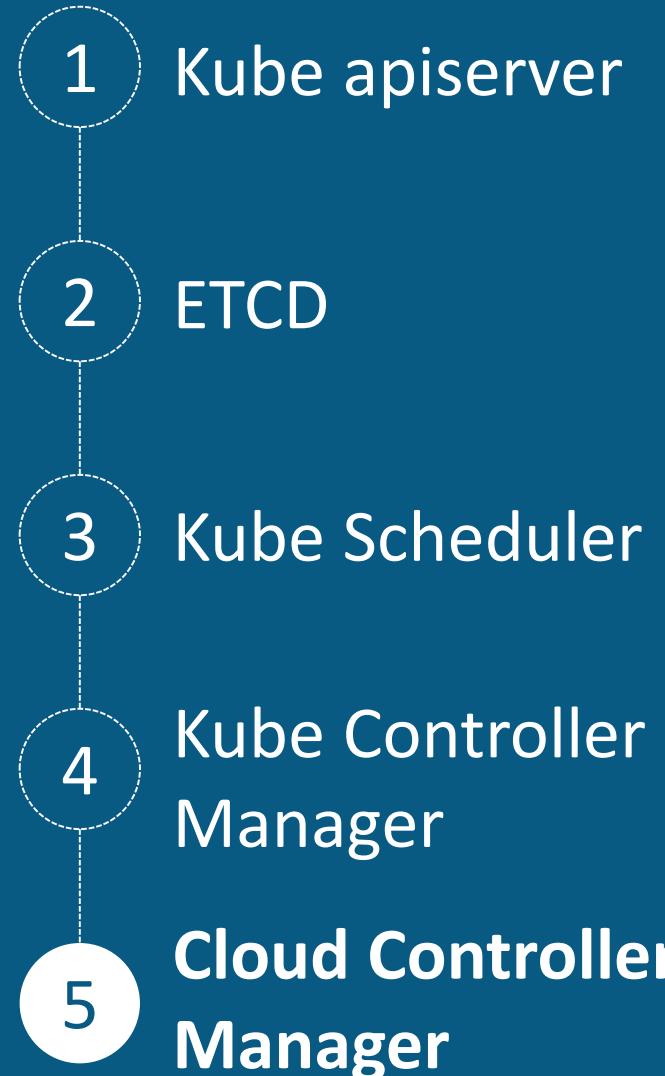


- There are multiple controller processes running on the master node, but are compiled together to run as a single process
- Each controllers has it's own responsibility and communicate with the API server to manage the endpoints
- Controller Manager is daemon that embed controllers
- Controller manager does namespace creation and garbage collection

# Controller runnings on master

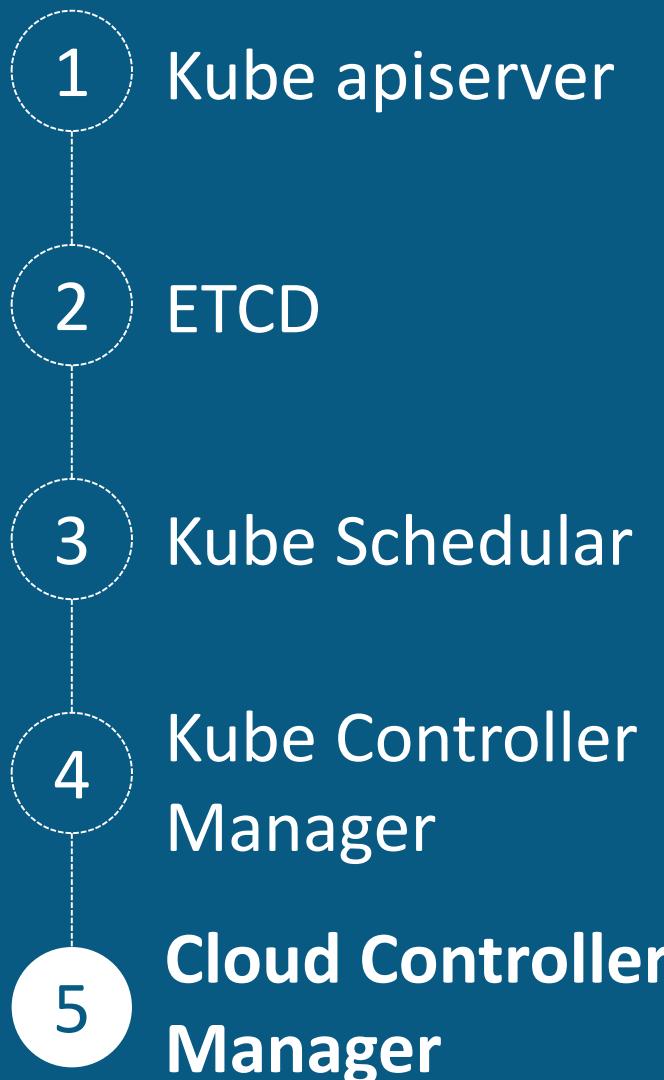


## Master Node Components



- Manages communication with the underlying cloud services
- It might be split out into several different container depending on which cloud platform you are running on
- It is responsible for persistent storage and network routing
- Cloud controller manager helps abstracting the cloud specific code from the core kubernetes specific code
- It enables cloud vendors and kubernetes code to be developed without any inter-dependency

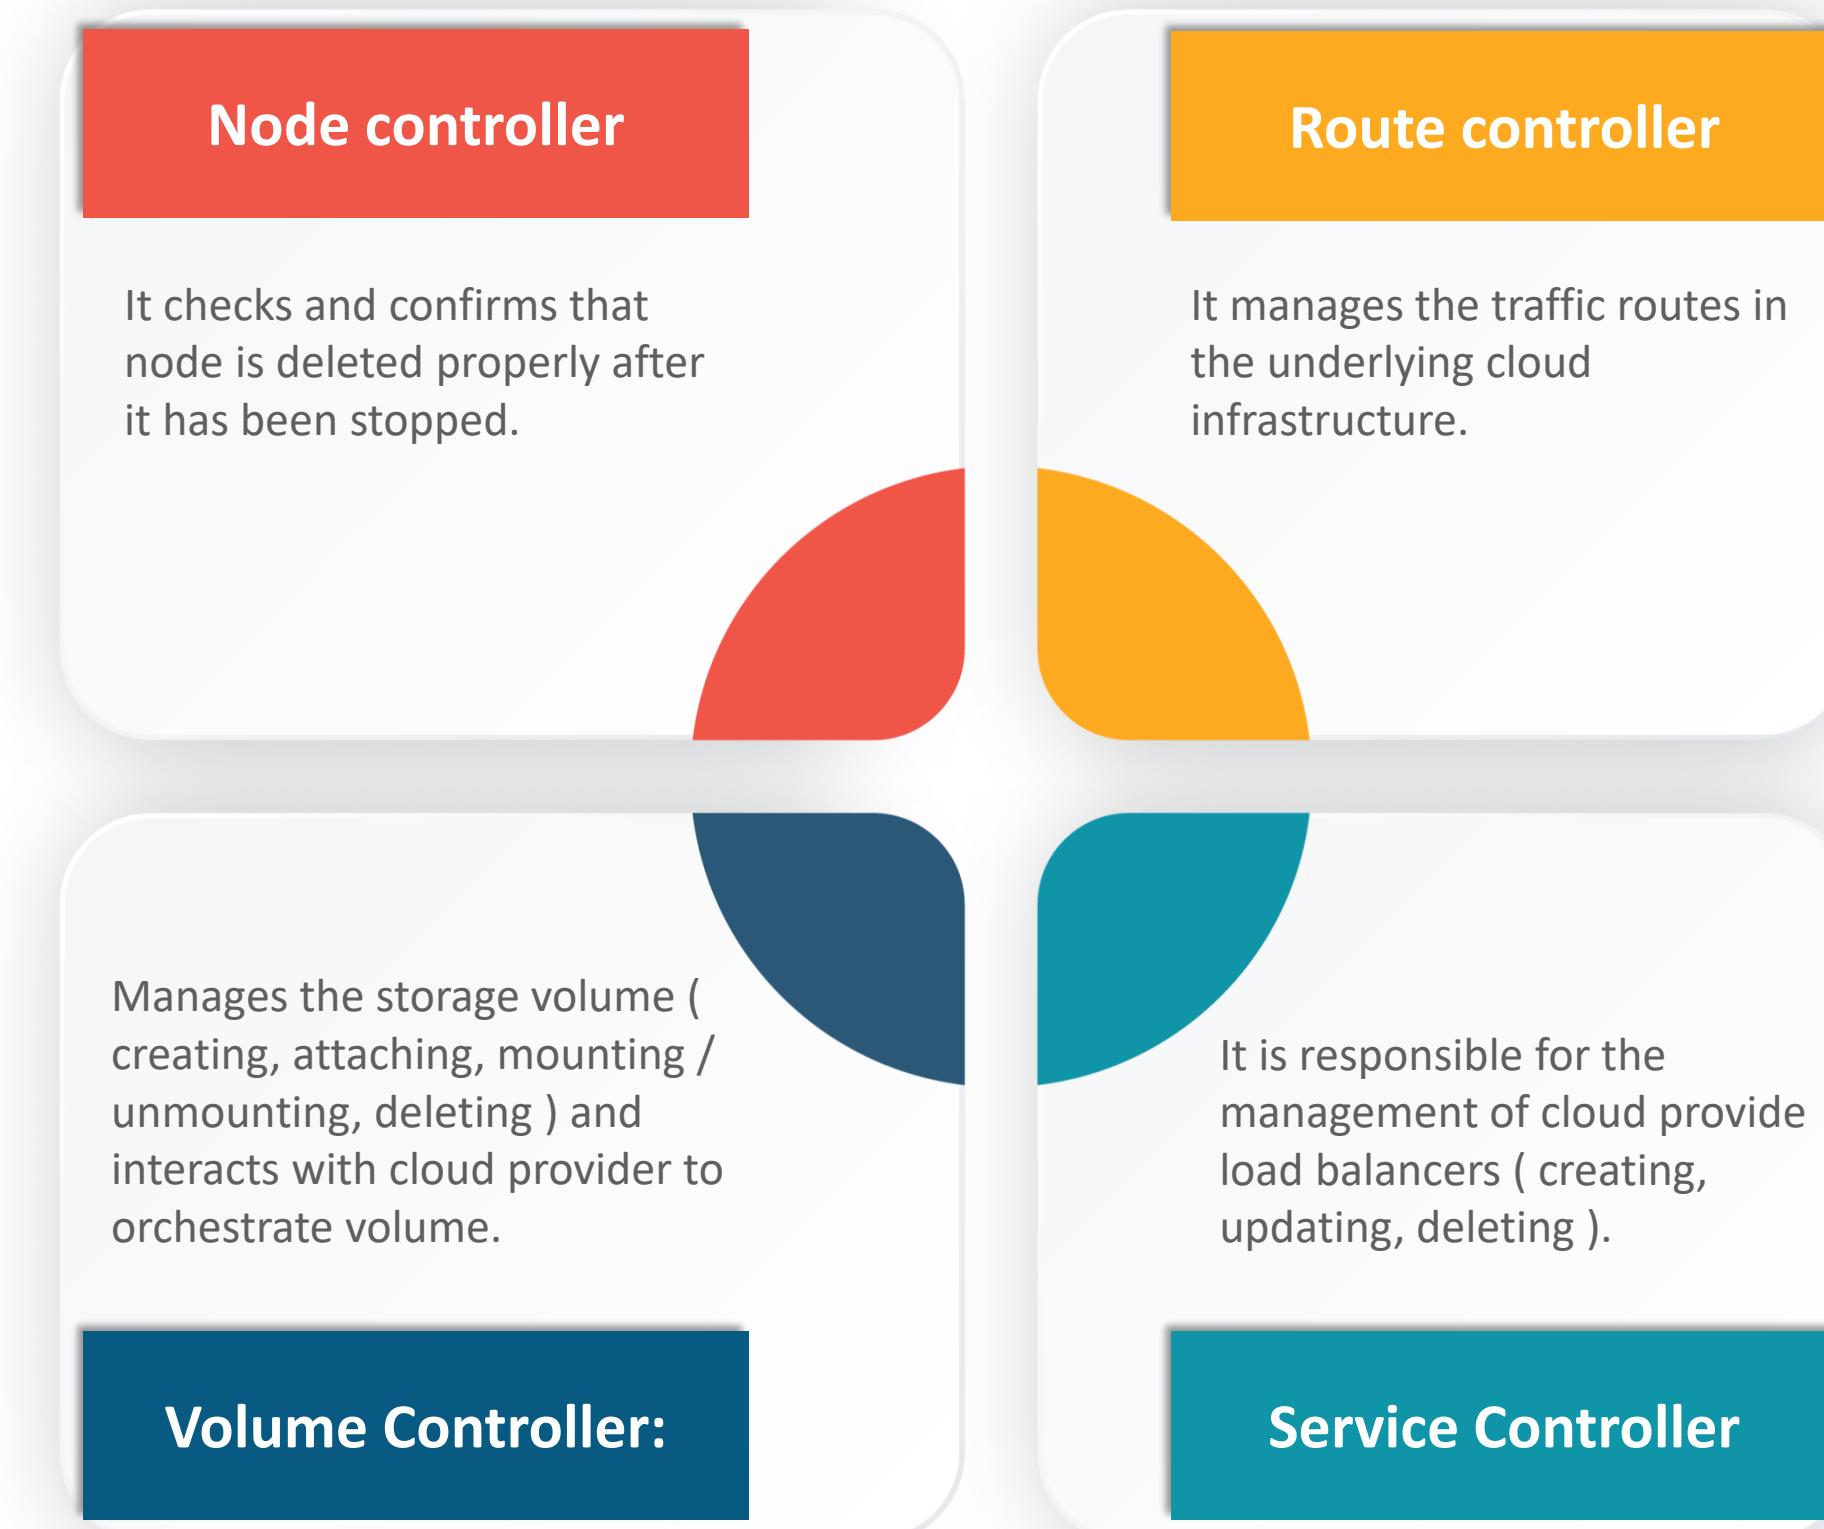
# Master Components of Kubernetes

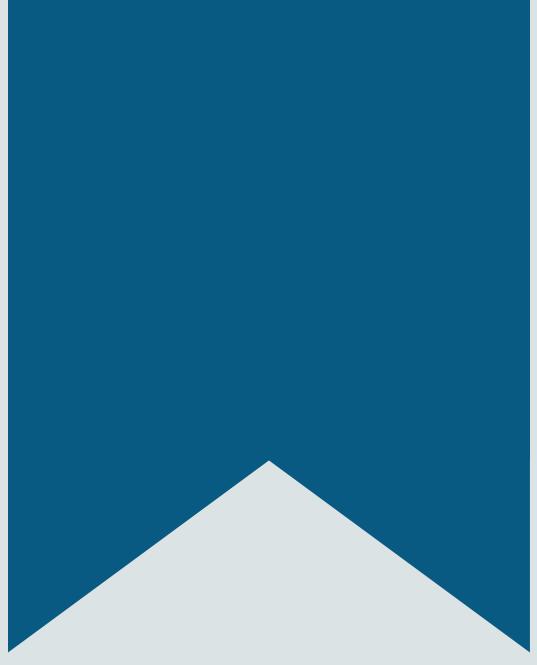


- Your cloud blocks the respective controller depending upon the use-case
- Cloud vendor develops their code and connects with kubernetes cloud-controller-manager while running the kubernetes

# Cloud-controller-manager Cont....:

- The cloud-controller has dependency on cloud service provider

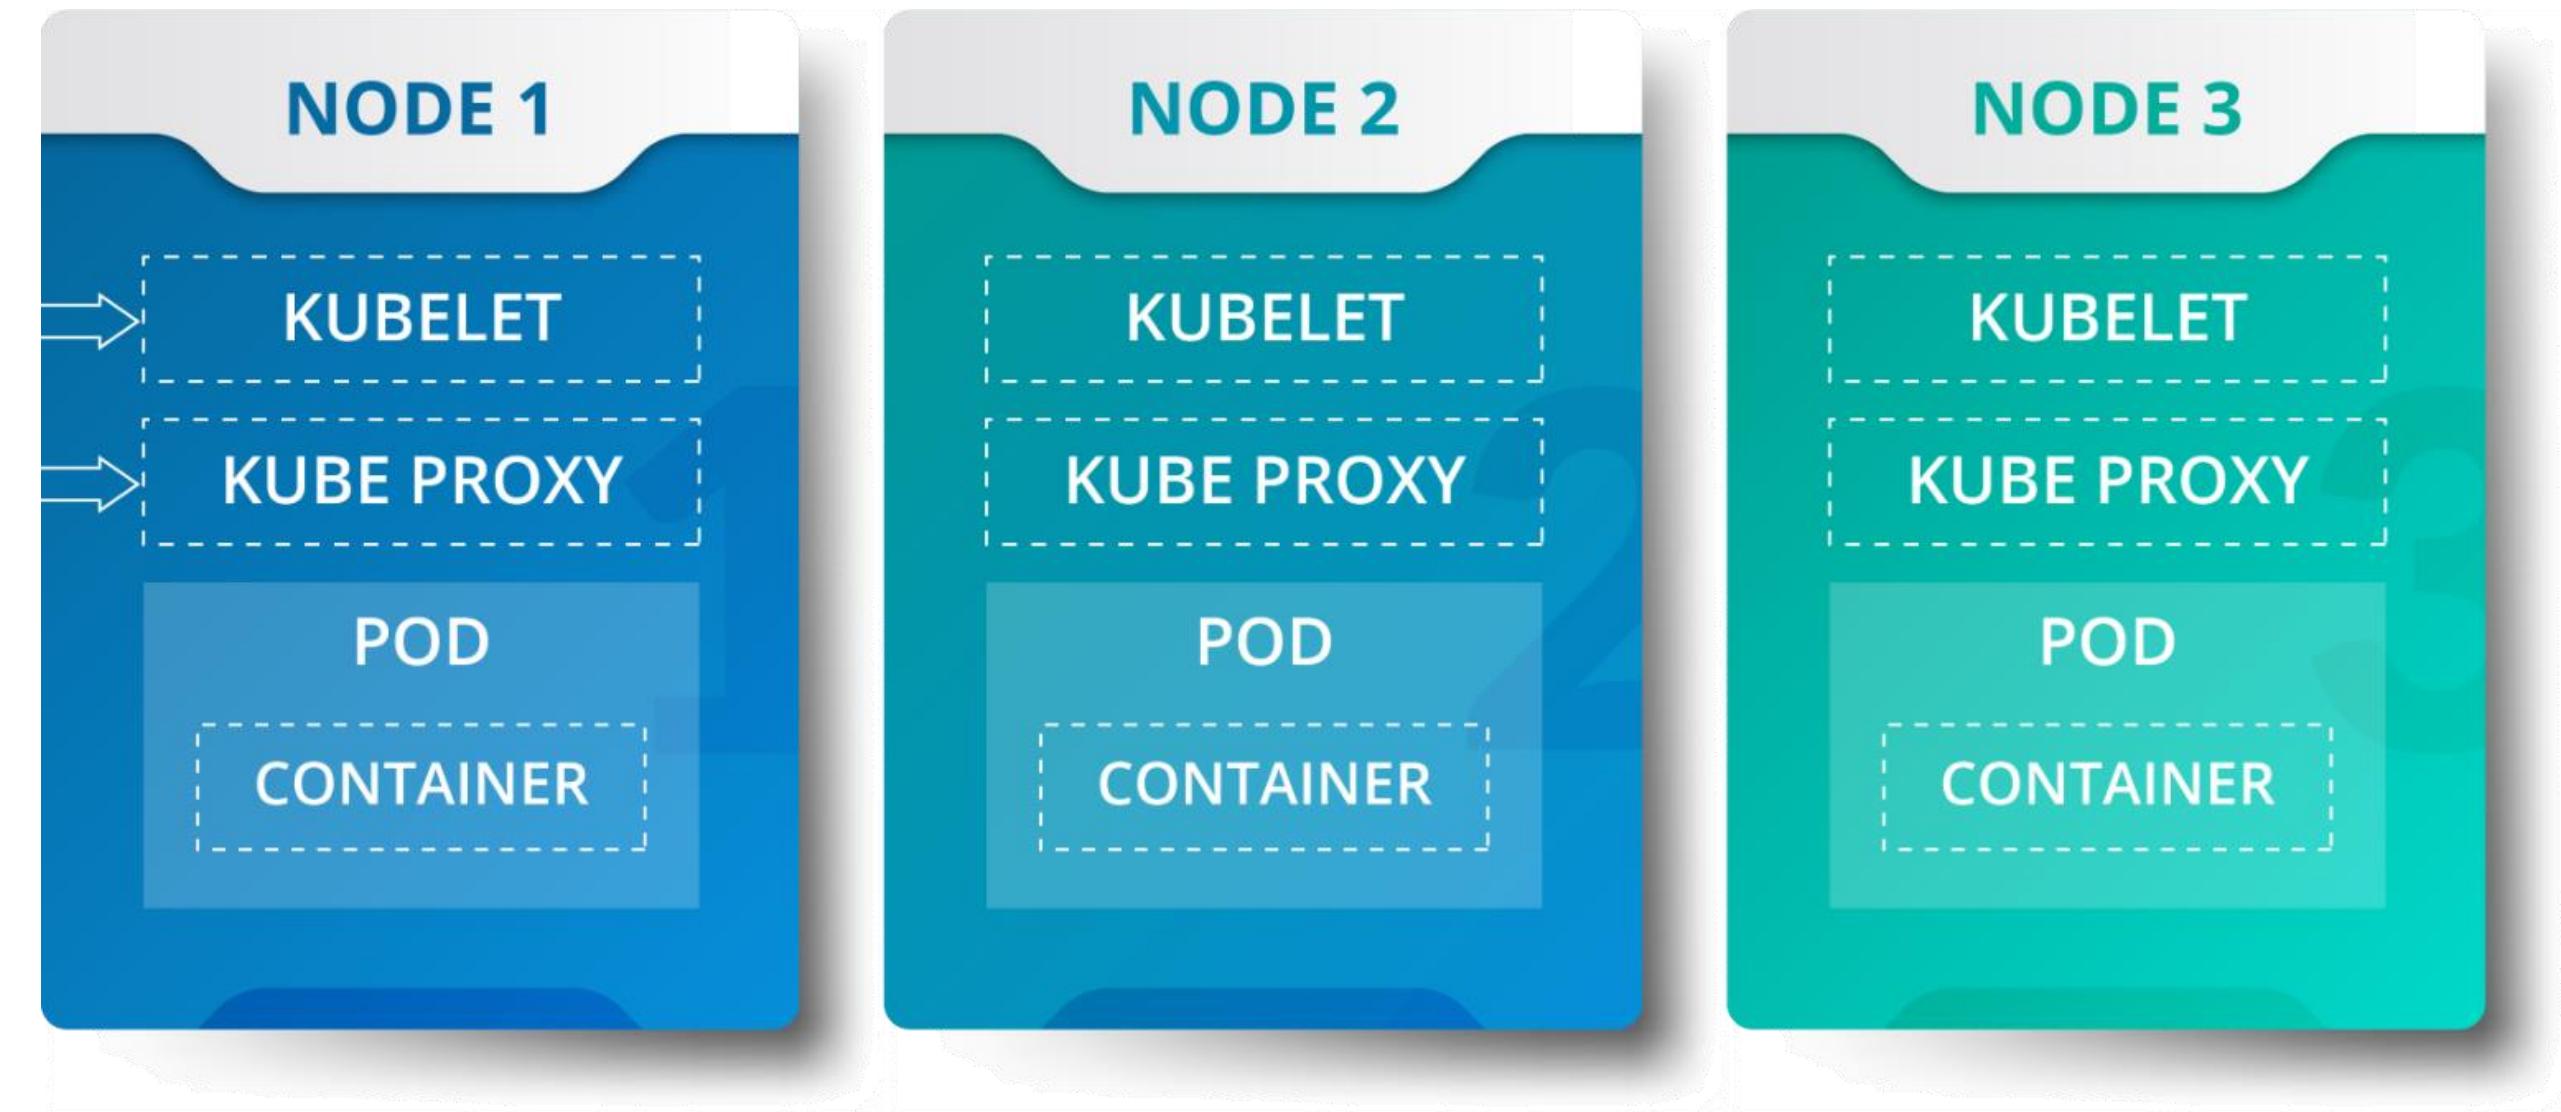




# Node components

# Node Components of Kubernetes

- kubelet
- kube-proxy
- kubectl
- Container (Docker)



# Node Components of Kubernetes

---

- Kubernetes nodes are the workhorses of the kubernetes cluster
- They were previously referred as Minions
- Customer workload or application runs on these nodes
- Multiple components run on each node:
  - Container runtime ( Docker )
  - kubelet
  - kube-proxy
  - kubectl

## Node Components of Kubernetes

- 1 Container Runtime
- 2 Kubelet
- 3 Kube-proxy
- 4 Kubectl

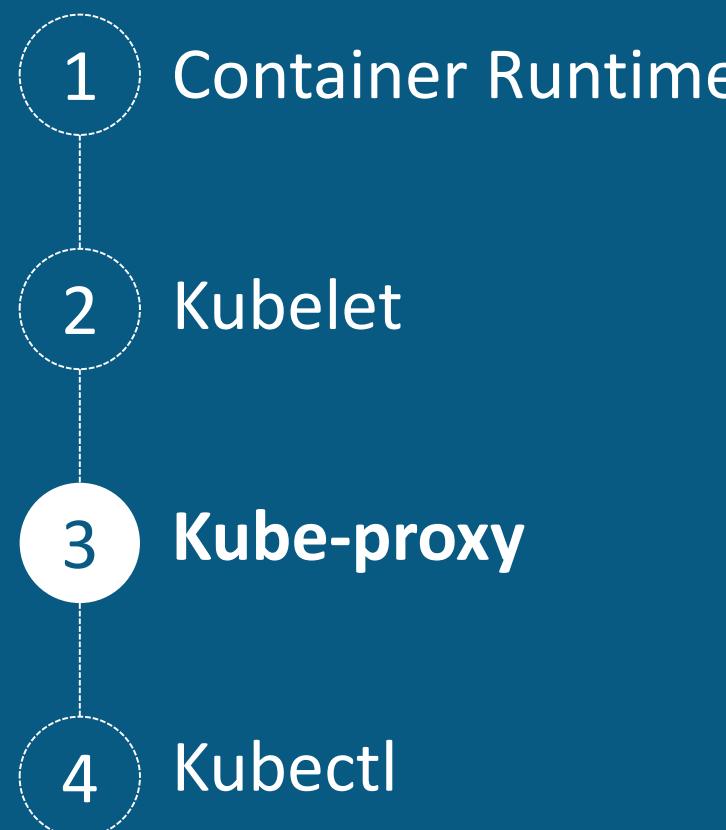
- **Container runtime** is the main software within a node which is responsible for running container
- **Docker**: Docker is just one but most popular container runtime which kubernetes supports
  - Other commonly used container runtime :
    - rkt
    - runc

## Node Components of Kubernetes

- 1 Container Runtime
- 2 Kubelet
- 3 Kube-proxy
- 4 Kubectl

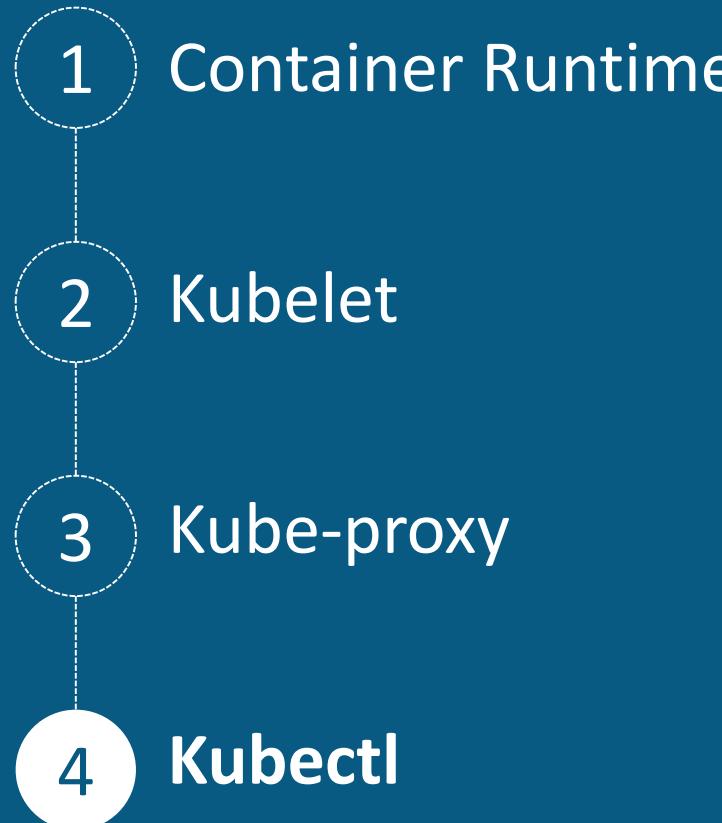
- Kubelet is an agent service which runs on each node and enables slave to communicate with master
- Kubelet works on PodSpec that are provided to it .  
PodSpec contains description of containers
- Kubelet ensures that containers described in PodSpec are Healthy and running
- It only manages the container that are created by Kubernetes master to which it belongs

## Node Components of Kubernetes



- Kubernetes kube-proxy runs on each node
- It can do simple TCP/UDP packet forwarding across backend network service
- It is a network proxy which reflects the services as configured in kubernetes API on each node
- Docker-link-compatible environment variables provides the cluster IPs and ports which are opened by proxy

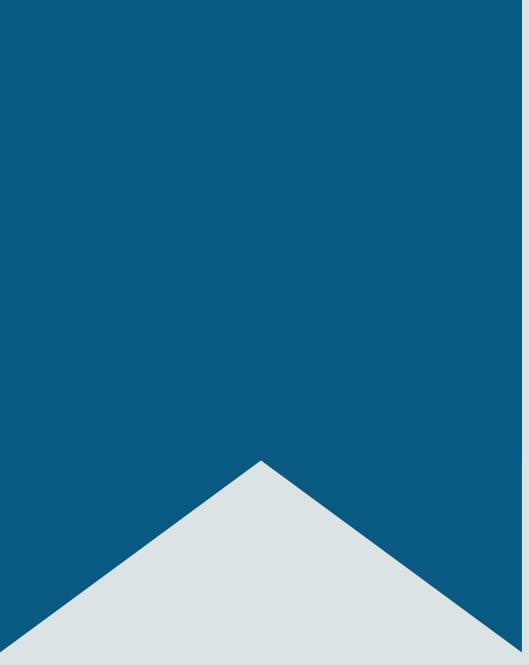
## Node Components of Kubernetes



- Kubectl is the interface/platform using which you can pass commands to the cluster
- It provides the CLI to run commands against the kubernetes cluster
- Kubectl command-line tool supports several different ways to create and manage kubernetes component



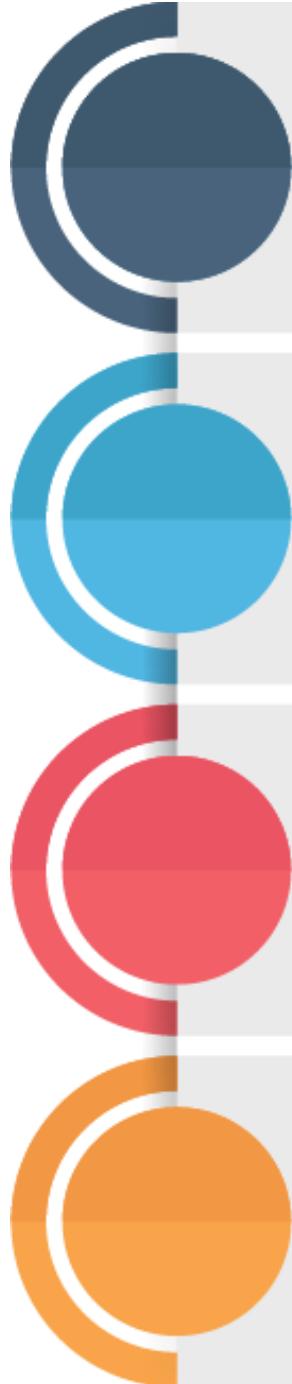
# **DEMO : Create a Kubernetes Cluster**



# Add-ons

# Addons: Cluster DNS

---



This is an optional item with regards to the functionality of the kubernetes cluster

For DNS, Kubernetes will create a DNS Pod and required services on the cluster

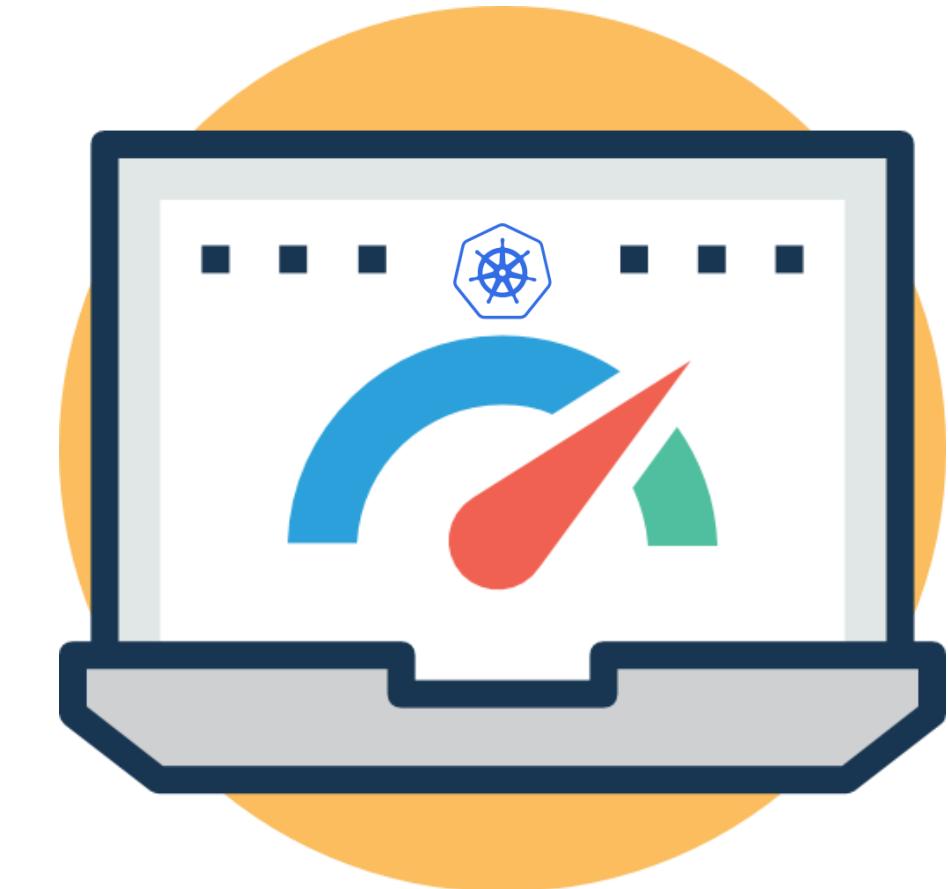
Through kubelets, it will update rest of the containers on the cluster with the DNS Service IP and do the name resolution

Every service in the kubernetes cluster is assigned a DNS name (including the DNS server itself)

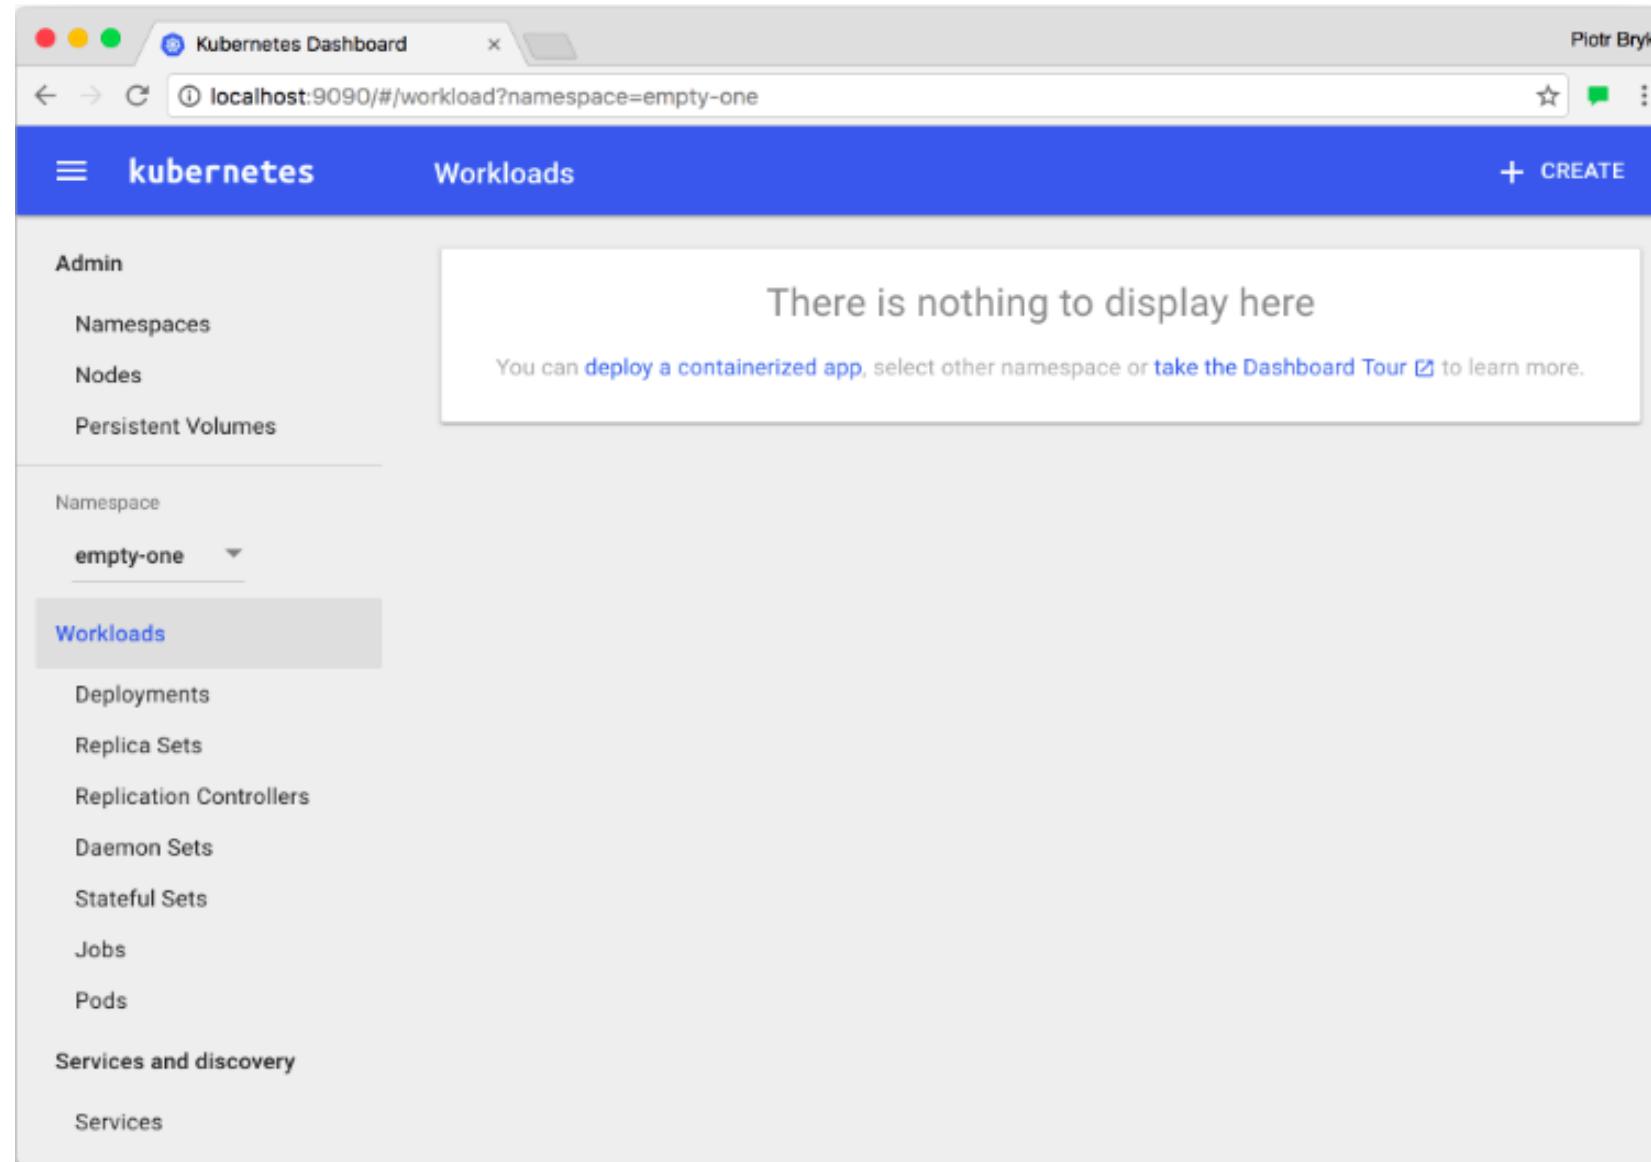
# Addons: Dashboard

---

- Kubernetes Dashboard provides the GUI capabilities to the cluster to operate and manage clusters
- It eases the job of admin / developer / IT manager
  - Helps to containerize application
  - Troubleshooting containerized application
  - Overview of all the applications running on the containerized application



# Addons: Kubernetes Dashboard



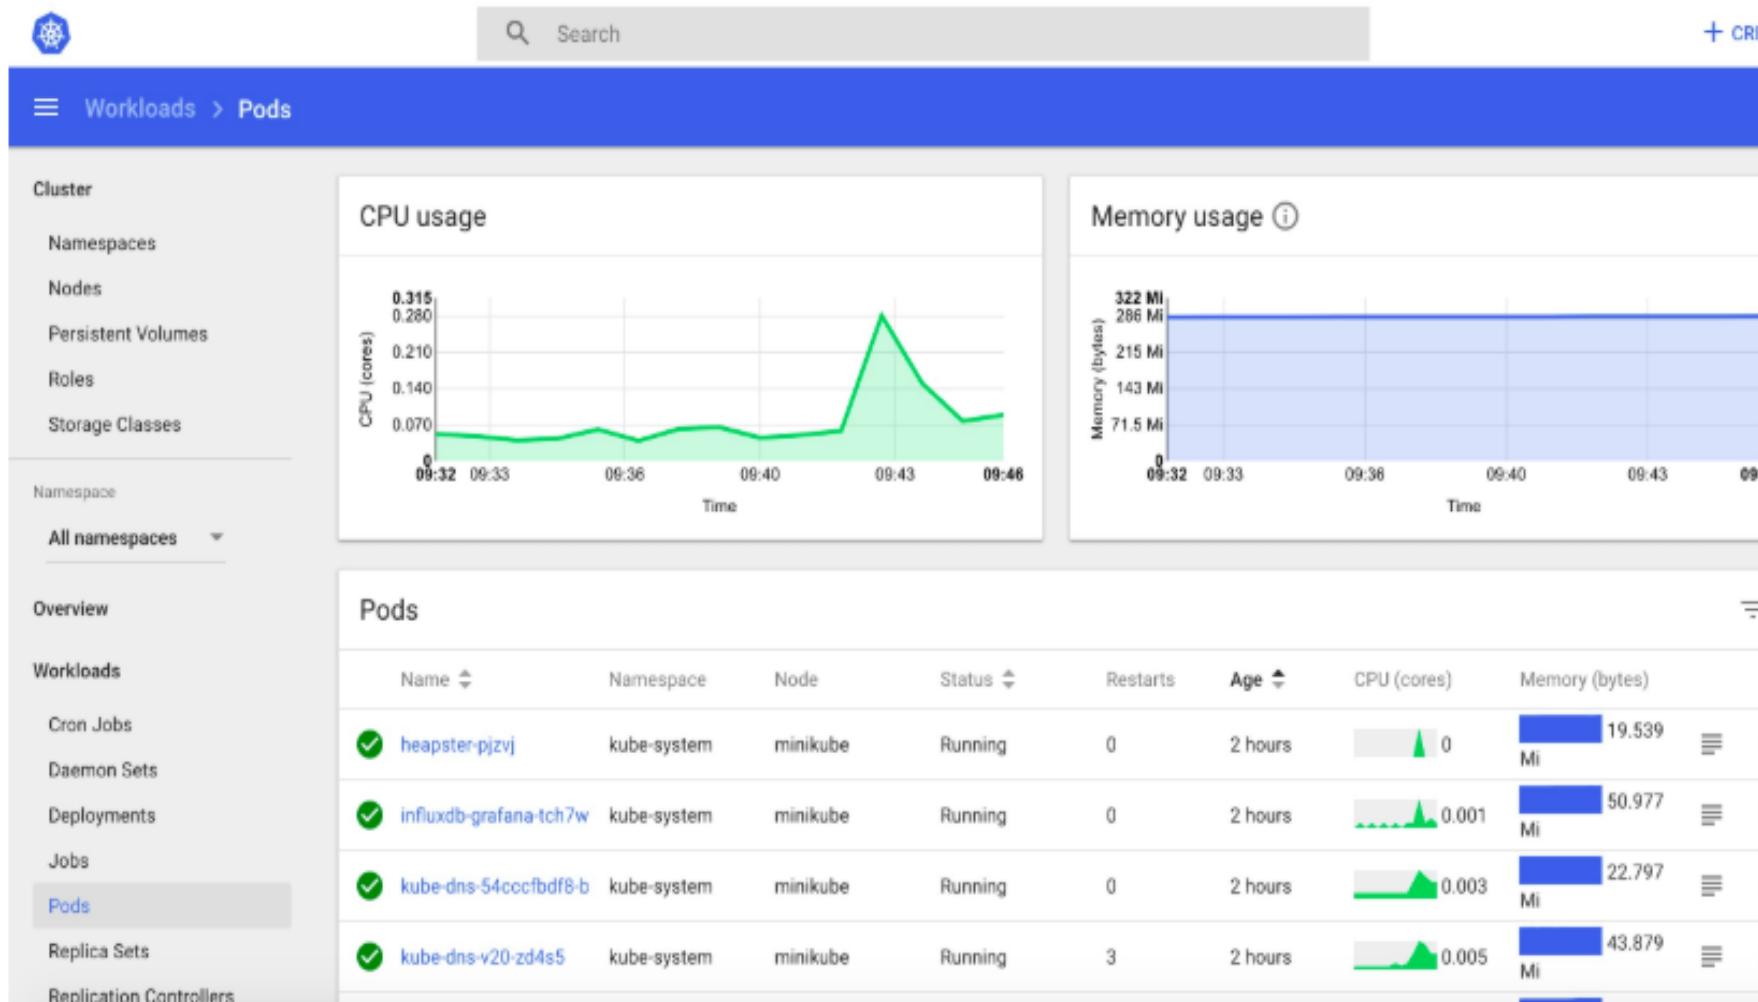
Ref : [Kubernetes Dashboard](#)

## For Developers :

- View to containerize application

A screenshot of the "Resource creation" dialog in the Kubernetes Dashboard. The title bar has a search icon and "+ CREATE". The main header says "Resource creation". On the left, a sidebar lists "Cluster" (with "Namespaces", "Nodes", "Persistent Volumes", "Roles", "Storage Classes"), "Namespace" (with "All namespaces" dropdown), "Overview", "Workloads" (with "Cron Jobs", "Daemon Sets", "Deployments", "Jobs", "Pods", "Replica Sets"), and "SHOW ADVANCED OPTIONS". The main content area has three tabs: "CREATE FROM TEXT INPUT", "CREATE FROM FILE" (selected), and "CREATE AN APP". Under "CREATE AN APP", fields include "App name\*" (with character count 0/24), "Container image\*", "Number of pods\*" (set to 1), and "Service\*" (set to "None"). Below these are "DEPLOY" and "CANCEL" buttons. To the right of the form, there are explanatory notes: "An 'app' label with this value will be added to the Deployment and Service that get deployed.", "Enter the URL of a public image on any registry, or a private image hosted on Docker Hub or Google Container Registry.", "A Deployment will be created to maintain the desired number of pods across your cluster.", and " Optionally, an internal or external Service can be defined to map an incoming Port to a target Port seen by the container.".

# Addons: Kubernetes Dashboard



Ref : [Kubernetes Dashboard](#)

## Admin View

- Understand the usage of the cluster
- Workload overview

The screenshot shows the 'Deployment' view for the 'review-app' deployment. It includes sections for 'New Replica Set', 'Old Replica Sets', and 'Horizontal Pod Autoscalers'. The 'Deployments' link in the sidebar is highlighted.

**New Replica Set:**

Name	Labels	Pods	Age	Images
review-app-57159...	app: review-app pod-template-h...	1 / 2	28 days	kubernetesdashb...

**Old Replica Sets:**

Name	Labels	Pods	Age	Images
review-app-43809...	app: review-app pod-template-h...	1 / 1	28 days	kubernetesdashb...

**Horizontal Pod Autoscalers:**

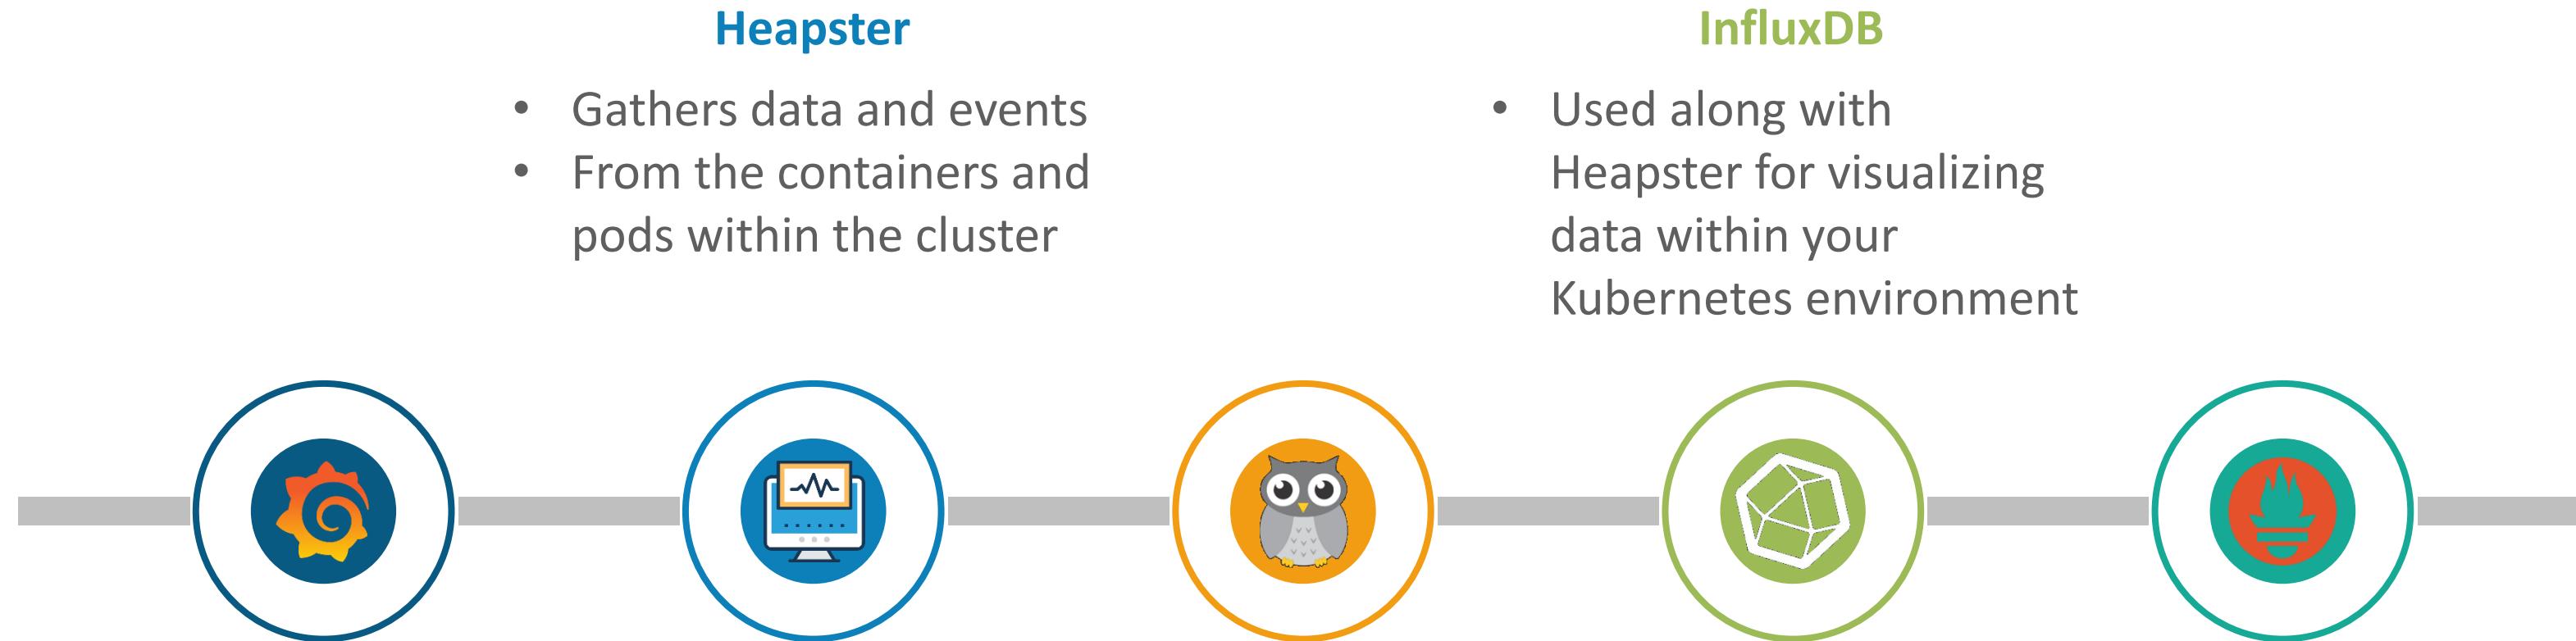
Name	Target CPU Utilization	Current CPU Utilization	Min Replicas	Max Replicas	Age
review-app	-	0%	2	8	10 days

# Addons: Container Resource Monitoring

---

- Monitoring of kubernetes cluster environment is much different from regular client-server architecture
- Kubernetes factored the management of the cluster by creating abstraction at different level like container, pods, services and whole cluster
- For users, it becomes very important to understand the performance of the application and resource utilization at all the different abstraction layer
- This approach is very important for large cluster and distributed resources

# Addons: Container Resource Monitoring Tools



## Heapster

- Gathers data and events
- From the containers and pods within the cluster

## InfluxDB

- Used along with Heapster for visualizing data within your Kubernetes environment

## Grafana

- A time-series database
- Stores data captured by all the Heapster pods

## CAdvisor

- Is a built-in tool in a kubelet that automatically discovers all the active containers monitors their performance and resource usage.

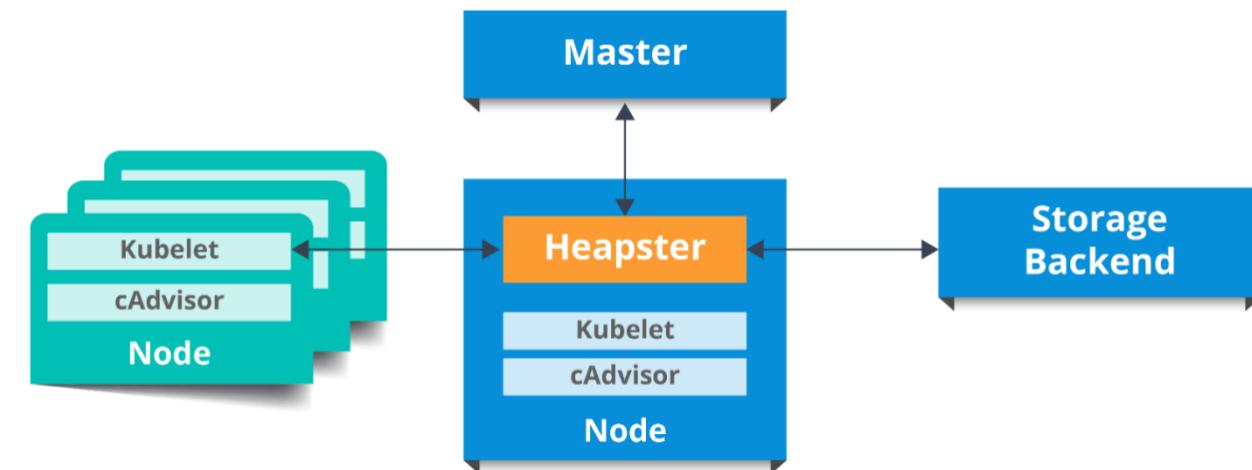
## Prometheus

- Project of CNCF ( Cloud Native Computing Foundation ), which provides a powerful querying, alerting and visualization capabilities

# Addons: Container Resource Monitoring Tools

---

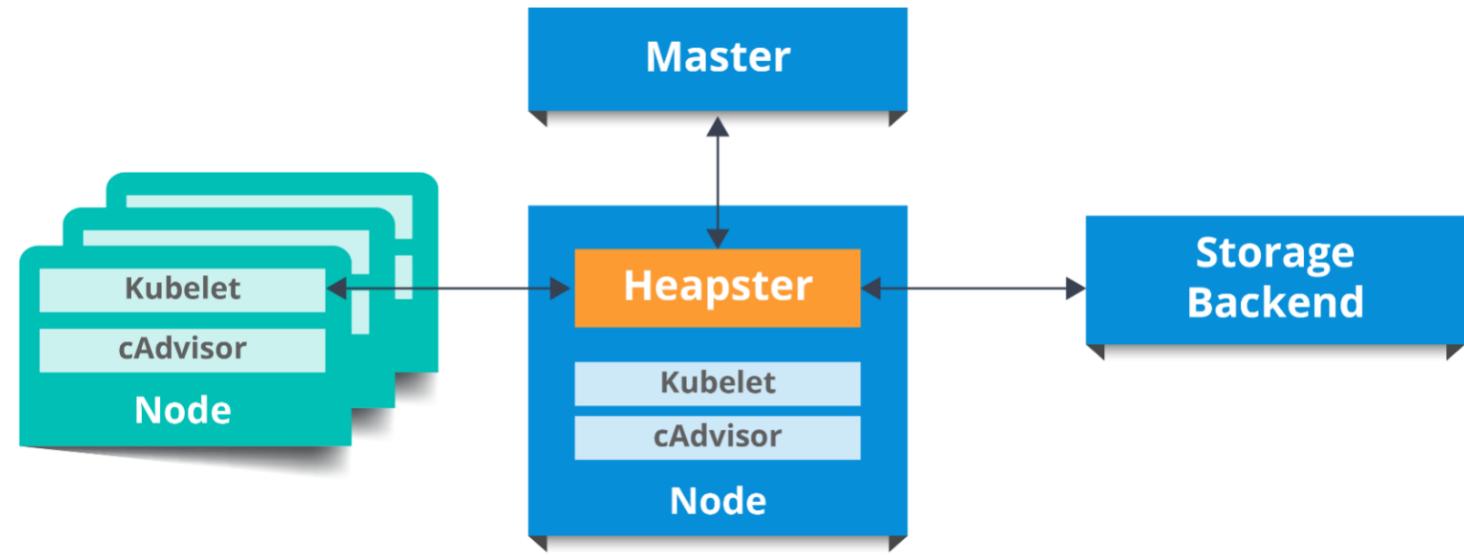
- On each node kubelet runs
- Kubelet fetches the data from cAdvisor
- This data is been passed on to Heapster
- Heapster classifies the data provided by kubelet and puts proper label to it
- This data is then pushed to the backend storage
- Visualization tools fetches the data from the backend storage and present it on the graphical dashboard



# Addons: Process Flow(Heapster)

Let's understand each components in more details:

- **Heapster:** It is a cluster wide aggregator of data provided by kubelet running on each node
- It is supported natively on kubernetes cluster
- It runs as a pod, just like any other pod in the kubernetes cluster
- It discovers all nodes in the cluster and queries usage information from the kubernetes nodes in the cluster, via on-machine kubernetes agent



# Addons: Process Flow(cAdvisor)

## cAdvisor:

- It is a resource usage and performance analysis agent
- It runs natively on docker containers and is a part of kubernetes binaries
- It runs on all the containers and collects :
  - CPU stats,
  - Memory stats,
  - Filesystem stats,
  - network usage statistics



# Installation of Tools – cAdvisor

---

```
$sudo docker run \
--volume=/:/rootfs:ro \
--volume=/var/run:/var/run:rw \
--volume=/sys:/sys:ro \
--volume=/var/lib/docker/:/var/lib/docker:ro \
--volume=/dev/disk/:/dev/disk:ro \
--publish=8080:8080 \
--detach=true \
--name=cadvisor \
google/cadvisor:latest
```

# Addons: Process Flow(InfluxDB)

## InfluxDB

- InfluxDB is used as a time-series database
- It is optimized for fast, highly available time-series queries used in operations, maintains and real-time analytics
- In kubernetes it runs as pod
- InfluxDB has no external dependencies
- It is accessed through sql-like language for querying built-in time-centric functions for data structures



# Installation of Tools – Heapster using InfluxDB

---

```
$ kubectl create -f deploy/kube-config/influxdb/  
$ kubectl create -f deploy/kube-config/rbac/heapster-rbac.yaml
```



YAML file can be taken from:

<https://github.com/kubernetes/heapster/tree/master/deploy/kube-config/influxdb>

# Addons: Process Flow(Grafana)

## Grafana

- InfluxDB works with Grafana
- Grafana provides graphical representation of Data processed by heapster and stored in influxDB
- Grafana is also an opensource tool, known for time-series analytics and visualization
- Grafana dashboard can easily be optimised and expanded



# Installation of Tools - Grafana

---

- Installation

```
$ docker run -i -p 3000:3000 grafana/grafana
```

- Configuration:

- sqlite3 database - /var/lib/grafana
- configuration files - /etc/grafana/

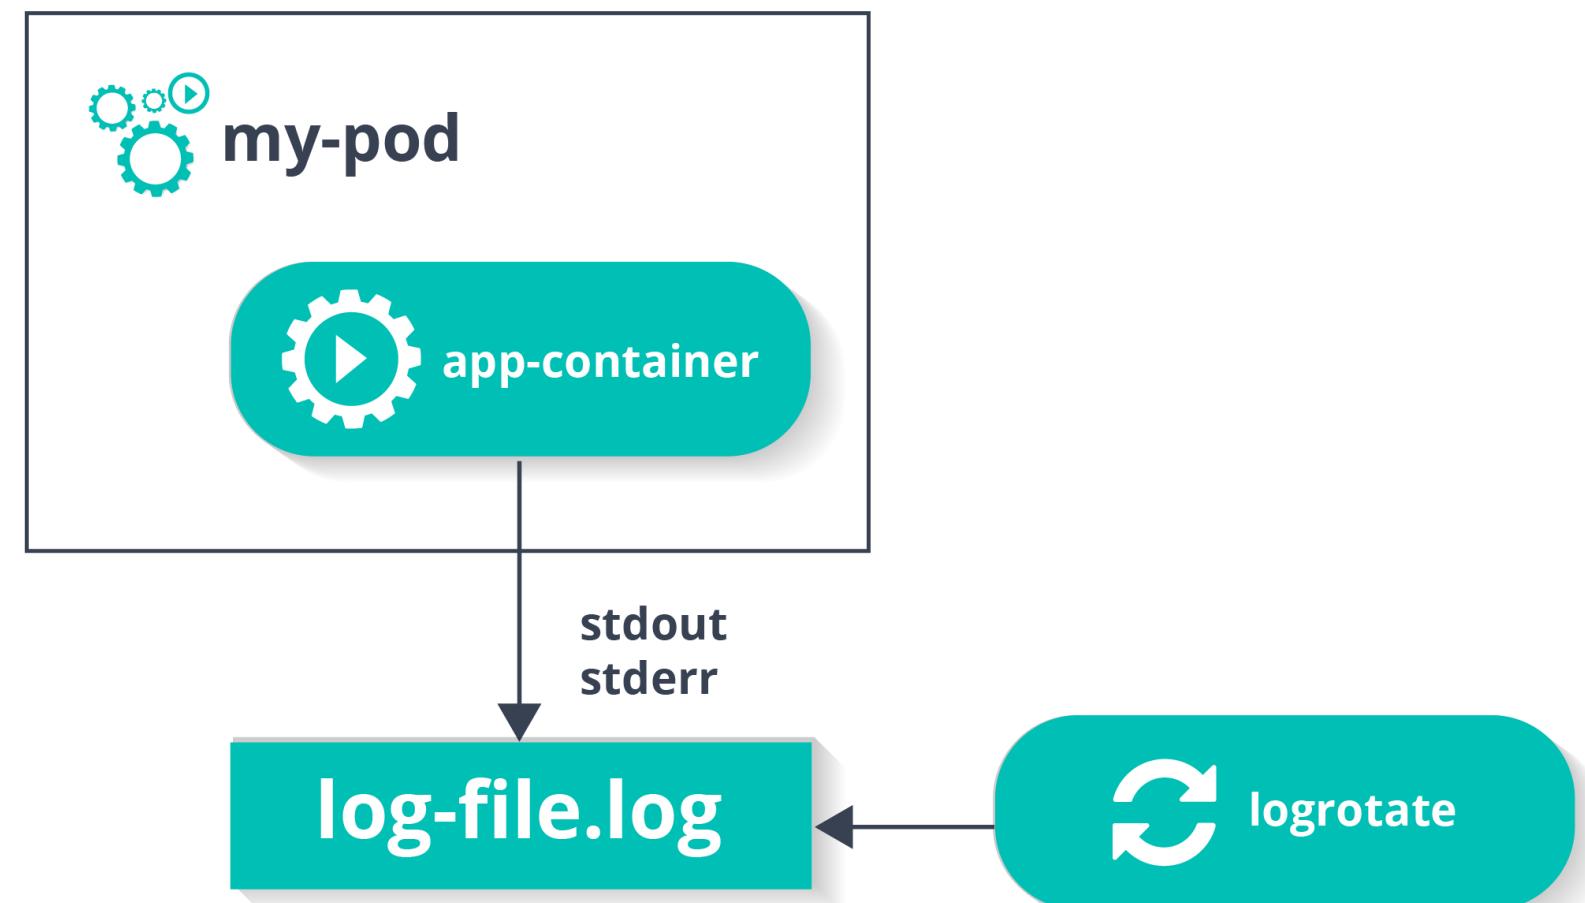
# Addons: Cluster level logging

---

- Cluster level logging requires additional backend to store
- By default, kubernetes does not have native storage for log data
- Cluster level logging assumes that there is a data-store inside or outside of kubernetes cluster
- All the containerized application writes to stdout
- Docker includes multiple logging mechanism
- Multiple logging mechanism helps to pull out the logs from each containers - These mechanisms are called logging driver

# Addons: Cluster level logging (Cont...)

- Kubernetes does not provide log rotation by default
- Additional log-rotation tools needs to be implemented to do so, avoiding logs to consume all the storage space



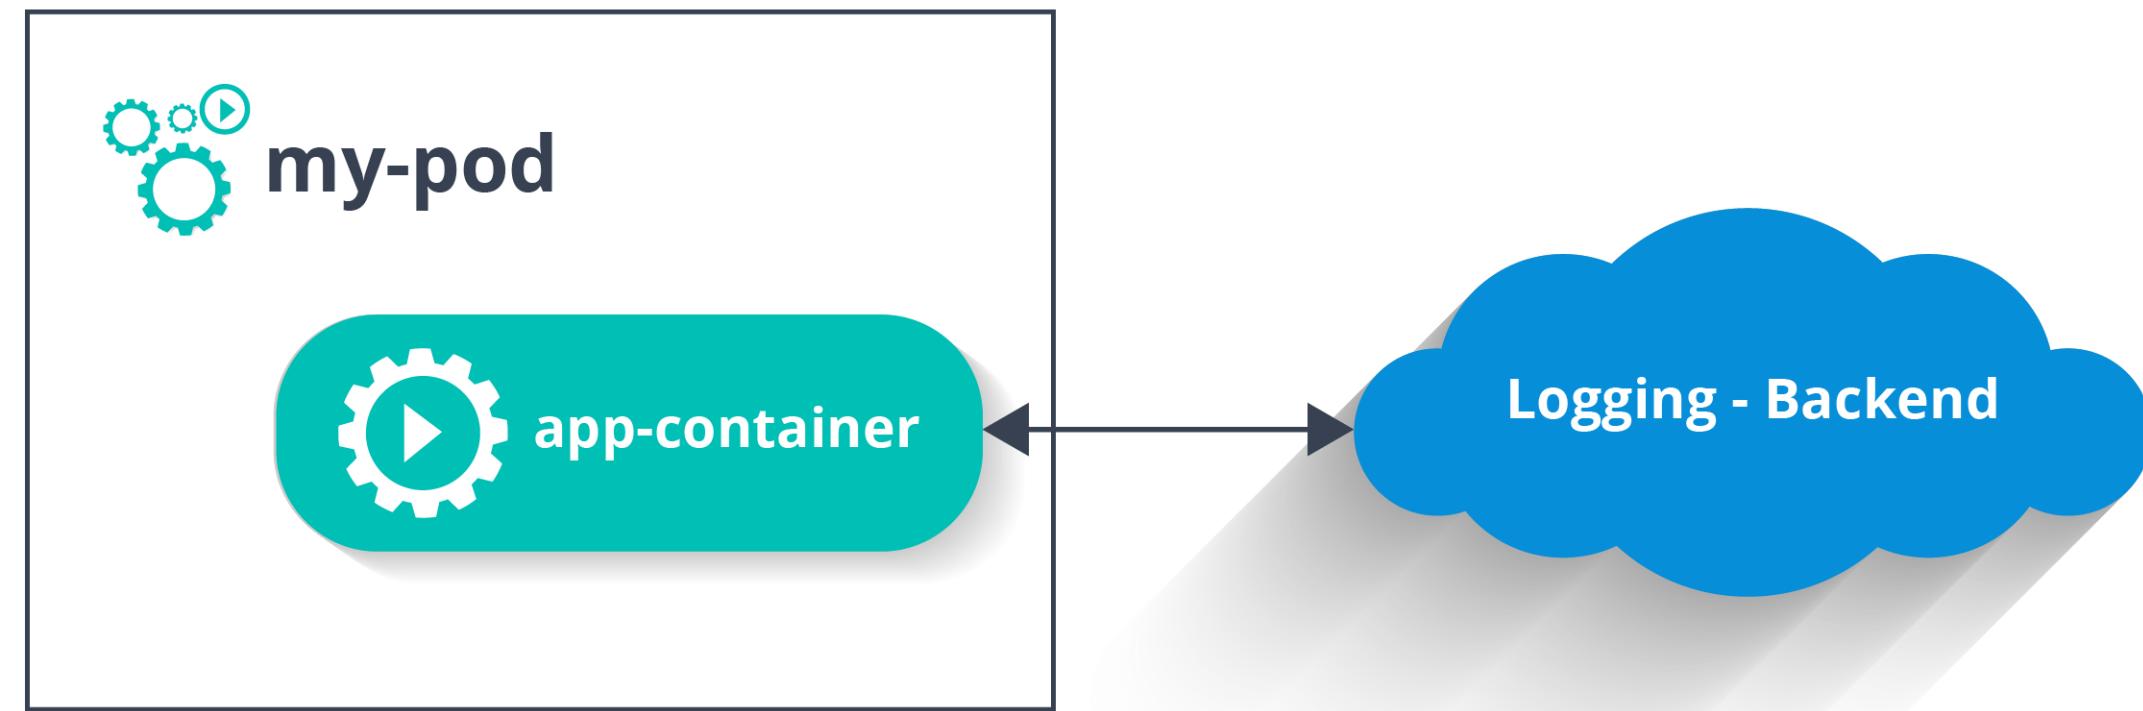
# Addons: Cluster level logging (Cont...)

---

- As mentioned earlier, kubernetes does not provide native solution for cluster-level logging. To handle this, here are few approaches. :
  - Configure mechanism to move the logs directly from application to a designated place like backend storage
  - Create another container next to main container ( also called sidcar) for logging purpose
  - Use an individual node-level agent to push the logs from individual to a designated place

Let's discuss each one of them in detail...

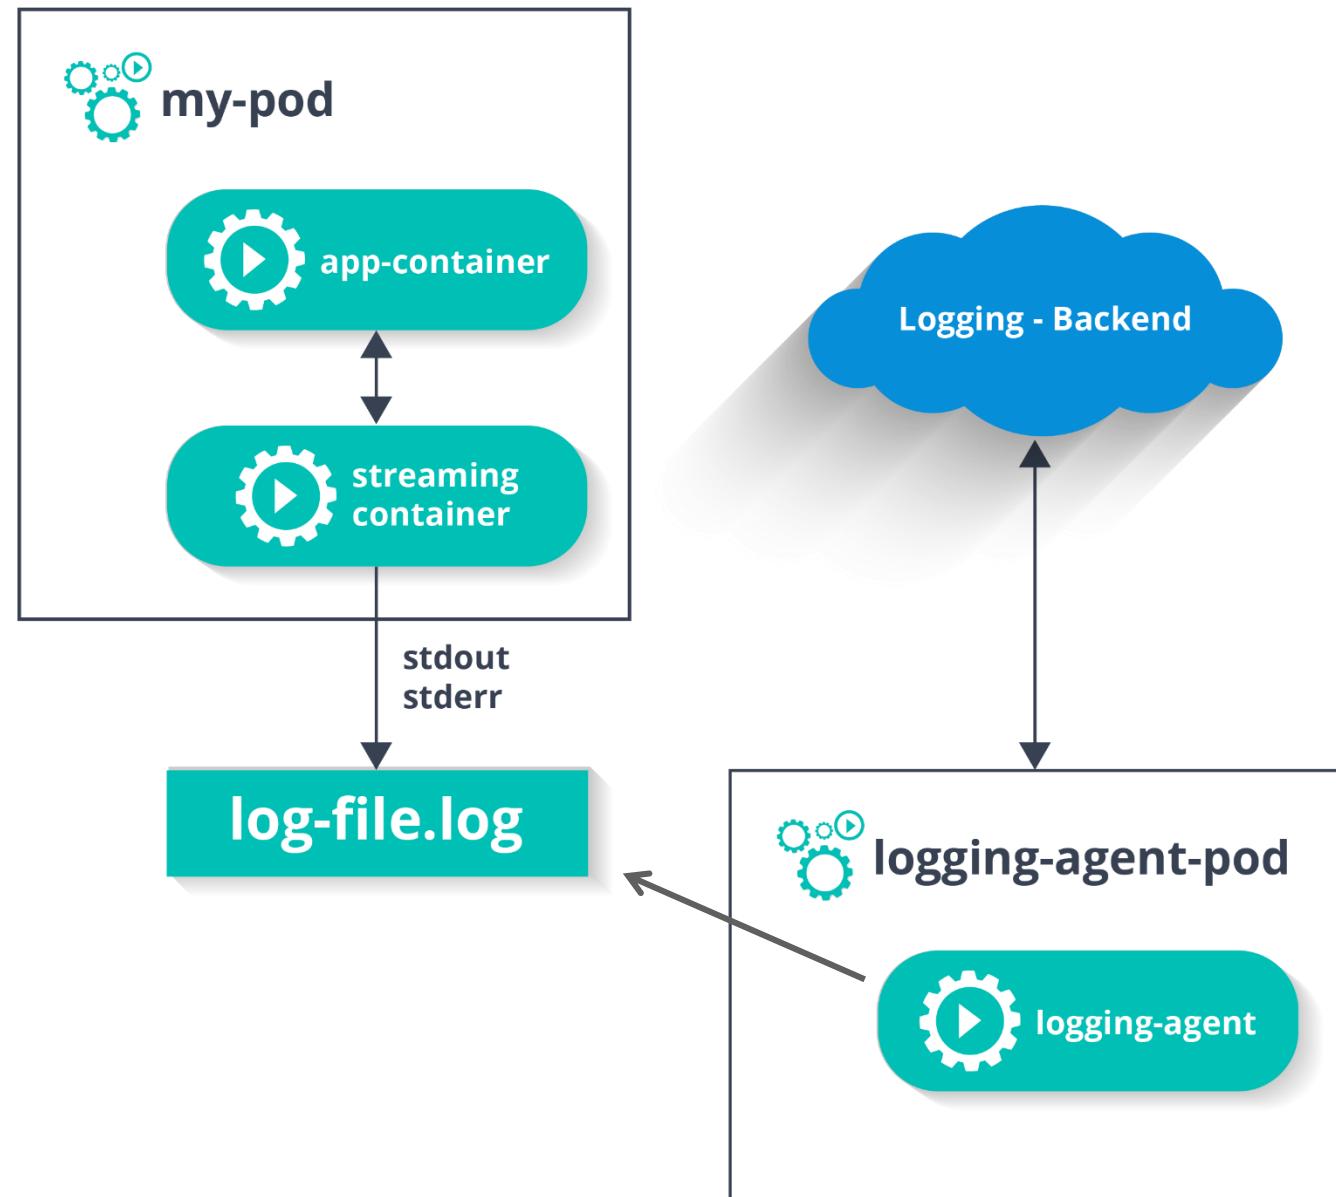
# Addons: Cluster level logging (Cont...)



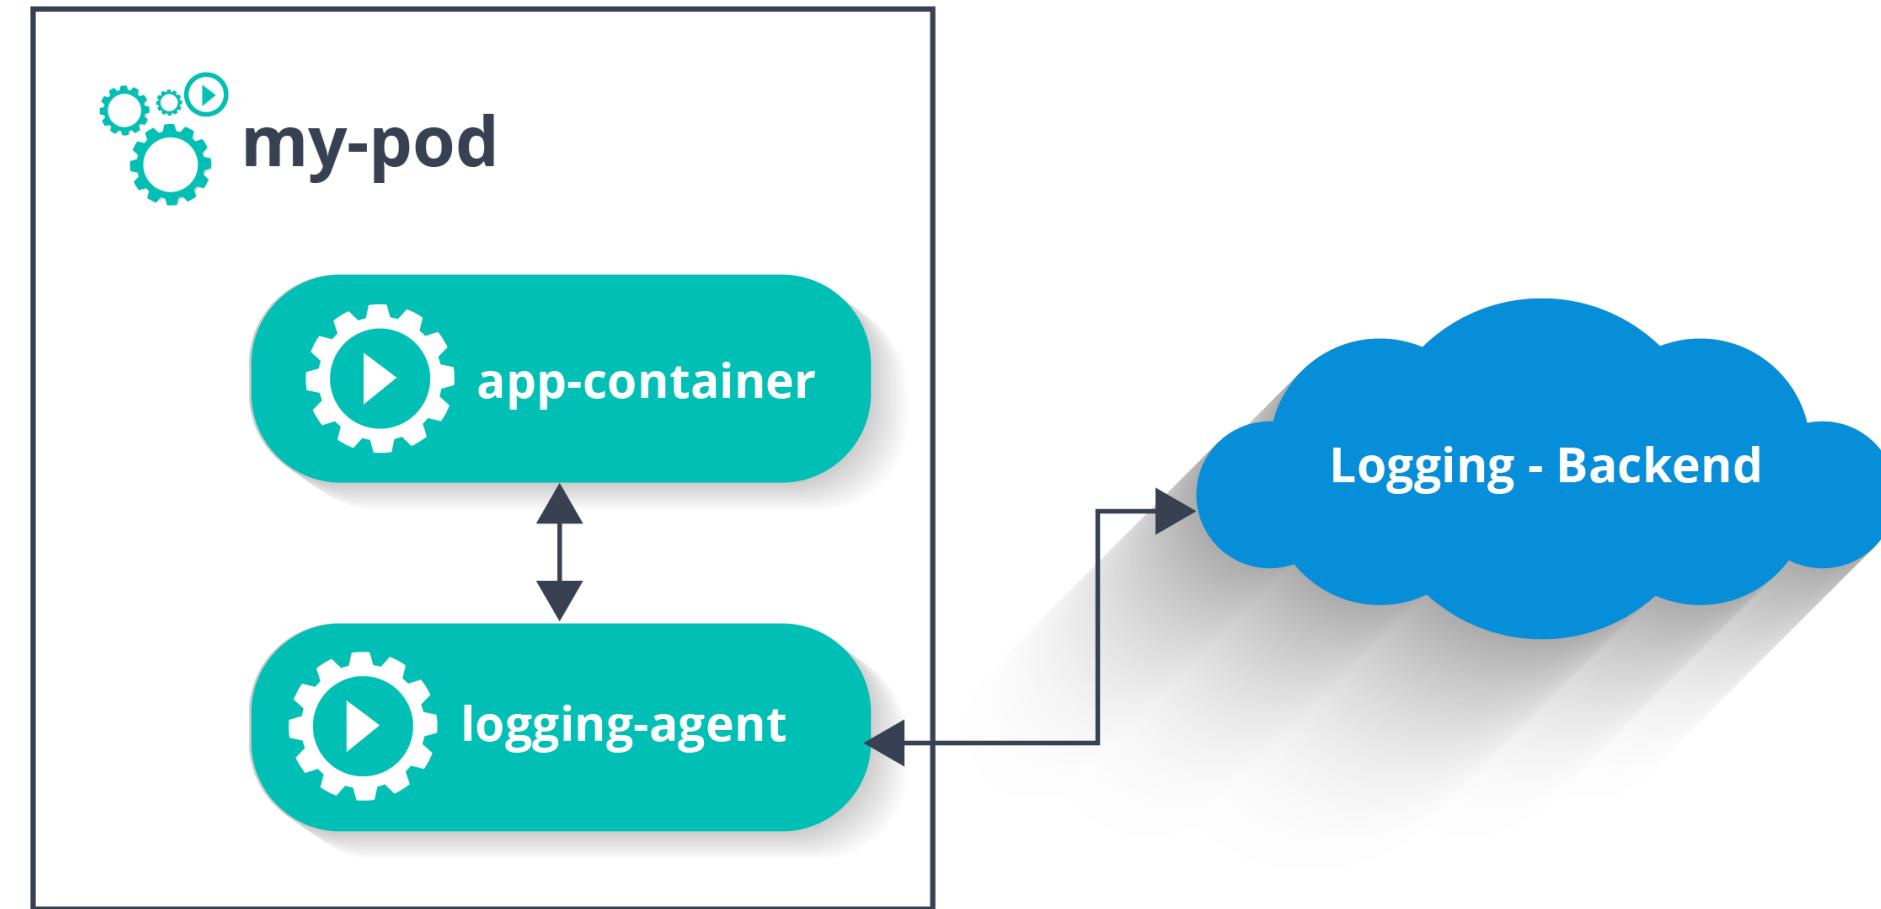
- Configure mechanism to move the logs directly from application to designated place like backend storage
- Configuration of this kind of logging is outside the scope of kubernetes and needs to be taken care by the application vendor / team

# Addons: Cluster level logging (Cont...)

- Create another container next to main container ( also called sidecar) for logging purpose
- In this approach, you can take the advantage of kubelet and logging agent
- Sidecar reads the logs from std-err and std-out from file and prints to its stream
- Another pods running with the logging agents on the cluster picks up the log
- This approach does not consume significant resource from the application container



# Addons: Cluster level logging (Cont...)



- This container runs the side-car container with logging agent running to it
- Drawback with this approach:
  - Side-car container consumes lot of resources
  - Logs can not be accessed using kubectl

# Kubeadm

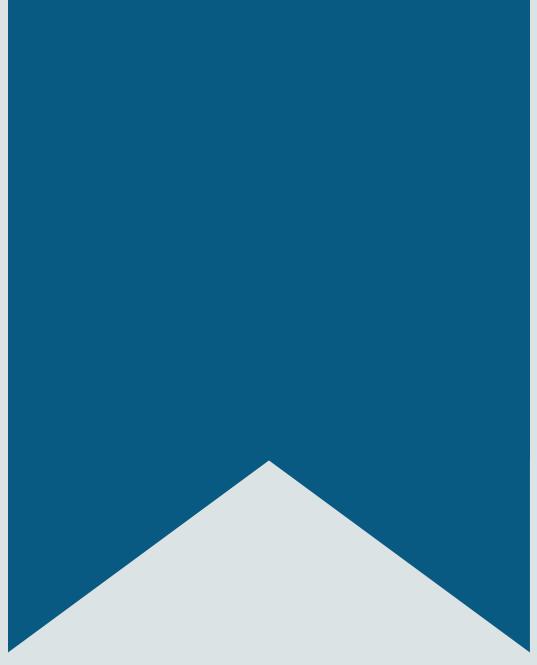
---

- For test and practise point of view, minikube was the first project which made the installation of kubernetes locally hasslefree
- Later, ubuntu introduced juju which allowed to deploy multi node kubernetes cluster in ubuntu
- Now, there is a new project from Mirantis which gives users the ability to create multi-node kubernetes clusters on any operation system where you can run docker
- This helps to understand how kubernetes clusters works , how they are created in production environment
- Kubeadm - kubeadm is a tool to bootstrap kubernetes cluster on existing infrastructure
- In this lab, we will setup the cluster on node-2

# Pre-requisite

---

- Install VM ( which you have done. )
- Install docker as you did in hands-on in Module-1
- In this lab, we will setup the cluster on node-2



# DEMO : Managing Cluster

# Get Cluster Details

- Check docker containers, pods

```
sudo kubectl get nodes  
sudo docker ps  
sudo kubectl get pods --all-namespaces  
Sudo kubectl cluster-info
```

```
[sachin@node-2:~/kubeadm-dind-cluster$ sudo kubectl cluster-info  
Kubernetes master is running at http://localhost:8080  
KubeDNS is running at http://localhost:8080/api/v1/namespaces/kube-system/services/kube-dns:dns/proxy  
To further debug and diagnose cluster problems, use 'kubectl cluster-info dump'.
```

```
[sachin@node-2:~/kubeadm-dind-cluster/fixed$  
[sachin@node-2:~/kubeadm-dind-cluster/fixed$ sudo kubectl get nodes  
NAME STATUS ROLES AGE VERSION  
kube-master Ready master 1h v1.8.11  
kube-node-1 NotReady <none> 1h v1.8.11  
kube-node-2 NotReady <none> 1h v1.8.11  
[sachin@node-2:~/kubeadm-dind-cluster/fixed$ sudo docker ps  
CONTAINER ID IMAGE COMMAND CREATED STATUS PORTS NAMES  
c35666f97646 mirantis/kubeadm-dind-cluster:v1.8 "/sbin/dind_init sys..." 24 minutes ago Up 24 minutes 8080/tcp kube-node-2  
a874fd72008a mirantis/kubeadm-dind-cluster:v1.8 "/sbin/dind_init sys..." 24 minutes ago Up 24 minutes 8080/tcp kube-node-1  
a3e076e14a84 mirantis/kubeadm-dind-cluster:v1.8 "/sbin/dind_init sys..." 24 minutes ago Up 24 minutes 127.0.0.1:8080->8080/tcp kube-master  
[sachin@node-2:~/kubeadm-dind-cluster/fixed$ sudo kubectl get pods --all-namespaces  
NAMESPACE NAME READY STATUS RESTARTS AGE  
kube-system etcd-kube-master 1/1 Running 2 1h  
kube-system kube-apiserver-kube-master 1/1 Running 2 1h  
kube-system kube-controller-manager-kube-master 1/1 Running 2 23m  
kube-system kube-dns-855bdc94cb-52dq5 3/3 Running 0 24m  
kube-system kube-proxy-kxrc4 1/1 Running 0 24m  
kube-system kube-proxy-mxcrf 1/1 Running 0 24m  
kube-system kube-proxy-svr8p 1/1 Running 0 24m  
kube-system kube-scheduler-kube-master 1/1 Running 2 1h  
kube-system kubernetes-dashboard-6b5bdctfc6-xc8w9 1/1 Running 0 24m  
[sachin@node-2:~/kubeadm-dind-cluster/fixed$
```

# Install KubeCtl

---

- Run these commands as root

```
Install kubelet kubeadm kubectl
#apt-get update && apt-get install -y apt-transport-https
curl
#curl -s https://packages.cloud.google.com/apt/doc/apt-
key.gpg | apt-key add -
#cat <<EOF >/etc/apt/sources.list.d/kubernetes.list
deb http://apt.kubernetes.io/ kubernetes-xenial main
EOF
#apt-get update
#apt-get install -y kubelet kubeadm kubectl
```

# Install KubeCtl

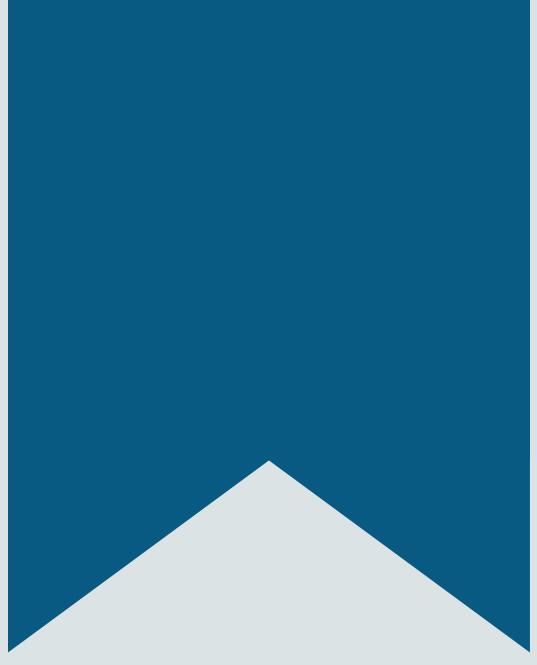
---

- Turn-off the swap

```
sachin@node-2:~/kubeadm-dind-cluster$ sudo swapoff -a  
sachin@node-2:~/kubeadm-dind-cluster$ sudo vi /etc/fstab  
sachin@node-2:~/kubeadm-dind-cluster$
```

- Set the permission

```
sachin@node-2:~/kubeadm-dind-cluster$ ls  
architecture.md  code-of-conduct.md  dind-cluster.sh  gce-setup.sh  OWNERS      SECURITY_CONTACTS  
AUTHORS.md      config.sh        Dockerfile       image          OWNERS_ALIASES  test.sh  
build          CONTRIBUTING.md   fixed           LICENSE        README.md  
sachin@node-2:~/kubeadm-dind-cluster$ cd fixed/  
sachin@node-2:~/kubeadm-dind-cluster/fixed$ ls  
dind-cluster-stable.sh  dind-cluster-v1.10.sh  dind-cluster-v1.8.sh  dind-cluster-v1.9.sh  
sachin@node-2:~/kubeadm-dind-cluster/fixed$ chmod +x dind-cluster-v1.8.sh
```



# DEMO: Create Kubernetes Dashboard

# Access The Dashboard

---

- Dashboard access weblink would be provided at the time of creation of the cluster

```
Warning: Stopping docker.service, but it can still be activated by:  
  docker.socket  
Warning: Stopping docker.service, but it can still be activated by:  
  docker.socket  
* Waiting for kube-proxy and the nodes  
.....[done]  
* Bringing up kube-dns and kubernetes-dashboard  
deployment "kube-dns" scaled  
deployment "kubernetes-dashboard" scaled  
.....[done]  
NAME        STATUS    ROLES      AGE       VERSION  
kube-master  Ready     master     3m        v1.8.11  
kube-node-1  Ready     <none>    2m        v1.8.11  
kube-node-2  Ready     <none>    2m        v1.8.11  
* Access dashboard at: http://localhost:8080/api/v1/namespaces/kube-system/services/kubernetes-dashboard/proxy
```

# Quiz

---

1. What is ETCD in kubernetes cluster ?
  - a. Key-value pair to store information for configuration
  - b. It is an English term to say “Et cetera” , meaning other things.
  - c. Selects the most suitable node to run the unscheduled pod
  - d. None of the above.

# Answers

1. What is ETCD in kubernetes cluster ?
  - a. Key-value pair to store information for container
  - b. It is an English term to say “Et cetera” , meaning other things.
  - c. Selects the most suitable node to run the unscheduled pod
  - d. None of the above.

**Answer A**



# Quiz

---

2. What are minions kubernetes cluster ?
  - a. They are cartoon character
  - b. They are work-horse / worker node of the kubernetes cluster
  - c. They are monitoring engine used widely in kubernetes
  - d. They are docker container service.

# Answers

2. What are minions kubernetes cluster ?
- a. They are cartoon character
  - b. They are work-horse / worker node of the kubernetes cluster
  - c. They are monitoring engine used widely in kubernetes
  - d. They are docker container service.

Answer B

# Quiz

3. Name controller-manager which has got dependency on cloud provider ?
- a. Node Controller : It checks and confirms that node is deleted properly after it has been stopped.
  - b. Route controller: It manager the traffic routes in the underlying cloud infrastructure.
  - c. Service Controller: It is responsible for the management of cloud provide load balancers ( creating, updating, deleting ).
  - d. Volume Controller: Manages the storage volume ( creating, attaching, mounting / unmounting, deleting ) and interacts with cloud provider to orchestrate volume
  - e. All of the above

# Answers

3. Name controller-manager which has got dependency on cloud provider ?
- a. Node Controller : It checks and confirms that node is deleted properly after it has been stopped.
  - b. Route controller: It manager the traffic routes in the underlying cloud infrastructure.
  - c. Service Controller: It is responsible for the management of cloud provide load balancers ( creating, updating, deleting ).
  - d. Volume Controller: Manages the storage volume ( creating, attaching, mounting / unmounting, deleting ) and interacts with cloud provider to orchestrate volume
  - e. All of the above

**Answer E:** All of the above



# Quiz

---

4. Kube-scheduler keeps track of resource utilization on all the compute nodes and ensure that workload is not scheduled on node which is already having scarcity of resources.
- a. True
  - b. False

# Answers

4. Kube-scheduler keeps track of resource utilization on all the compute nodes and ensure that workload is not scheduled on node which is already having scarcity of resources.
- a. True
  - b. False

**Answer A**



# Quiz

---

## 5. Map the definition correctly

- a. Heapster: 1) Comes as a built-in tool into kubelet and automatically discover all the active containers.  
It focuses on container level performance and resource usage.
- b. InfluxDB : 2) A time-series database that stores the data captured by all the Heapster pods
- c. Grafana: 3) Used along with Heapster for visualizing data within your Kubernetes environment
- d. CAdvisor: 4) It gathers data and events from the containers and pods within the cluster

# Answers

5. Map the definition correctly
- a. Heapster: 1) Comes as a built-in tool into kubelet and automatically discover all the active containers.  
It focuses on container level performance and resource usage.
  - b. InfluxDB : 2) A time-series database that stores the data captured by all the Heapster pods
  - c. Grafana: 3) Used along with Heapster for visualizing data within your Kubernetes environment
  - d. CAdvisor: 4) It gathers data and events from the containers and pods within the cluster

**Answer :**

A-4 , b-2, c-3, d-1



# Quiz

---

6. What is prometheus ?

# Answers

## 6. What is prometheus ?

**Answer :** Project of CNCF ( Cloud Native Computing Foundation ), which provides a powerful querying, alerting and visualization capabilities



# Quiz

---

7. Is heapster mandatory component for docker engine to work properly?

- a. Yes
- b. No

# Answers

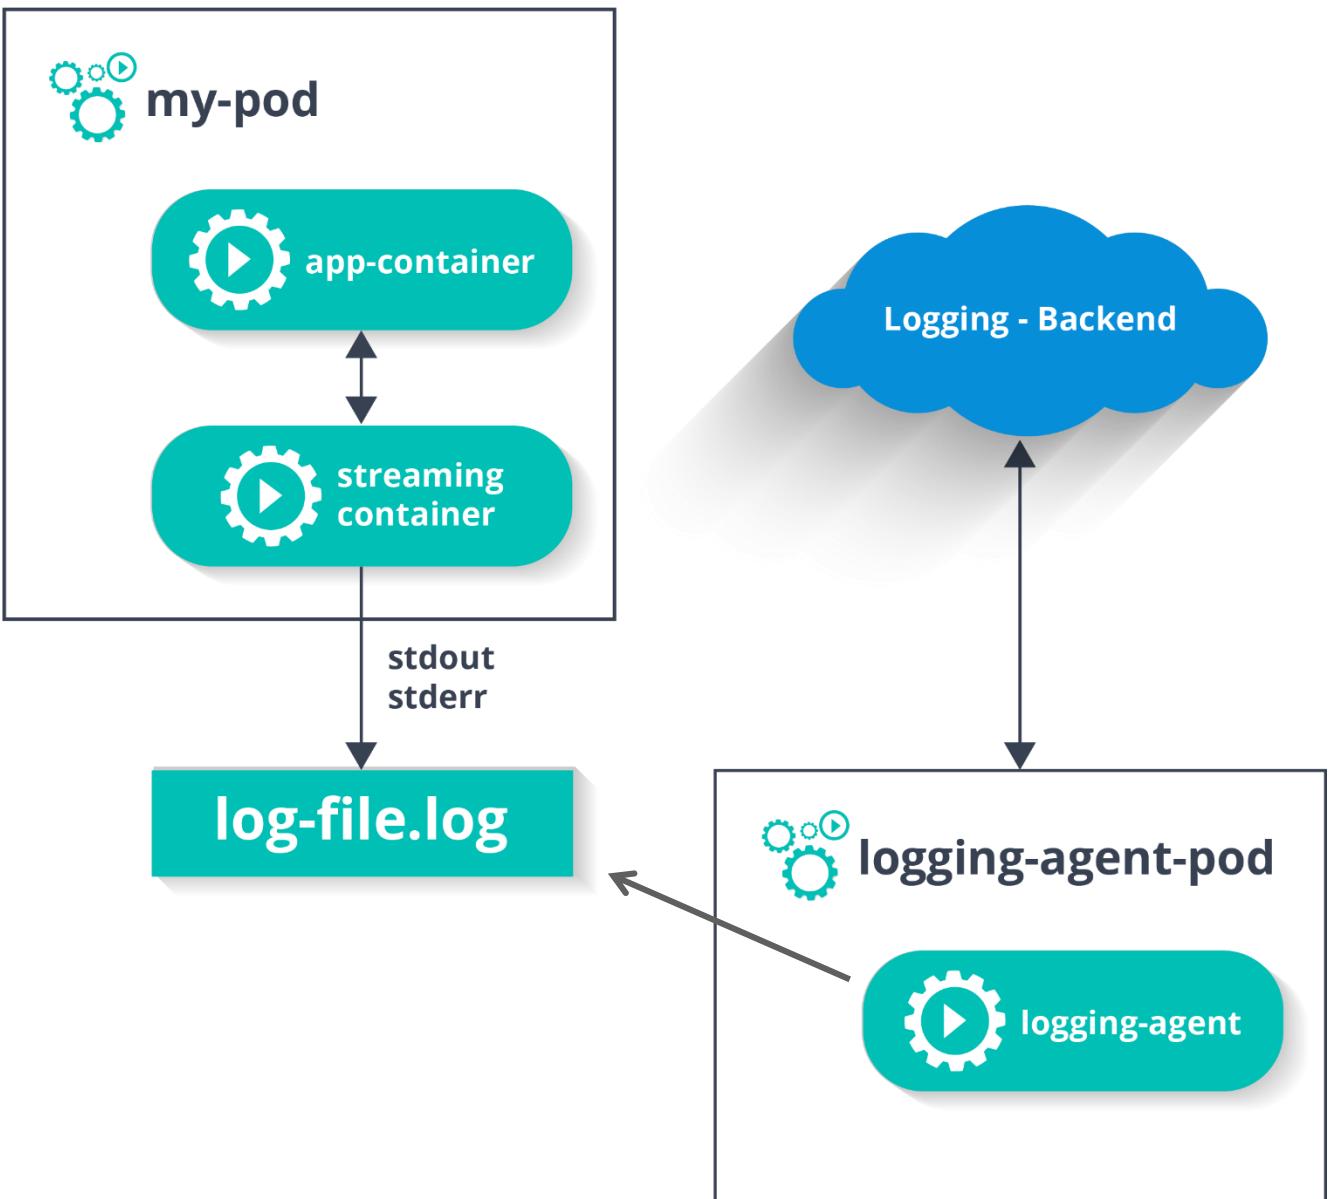
7. Is heapster mandatory component for docker engine to work properly?

- a. Yes
- b. No

**Answer B:** No, It is a cluster wide aggregator of data provided by kubelet running on each node. It is used for resource monitoring. Nothing to do with basic functionality of kubernetes or docker engine.

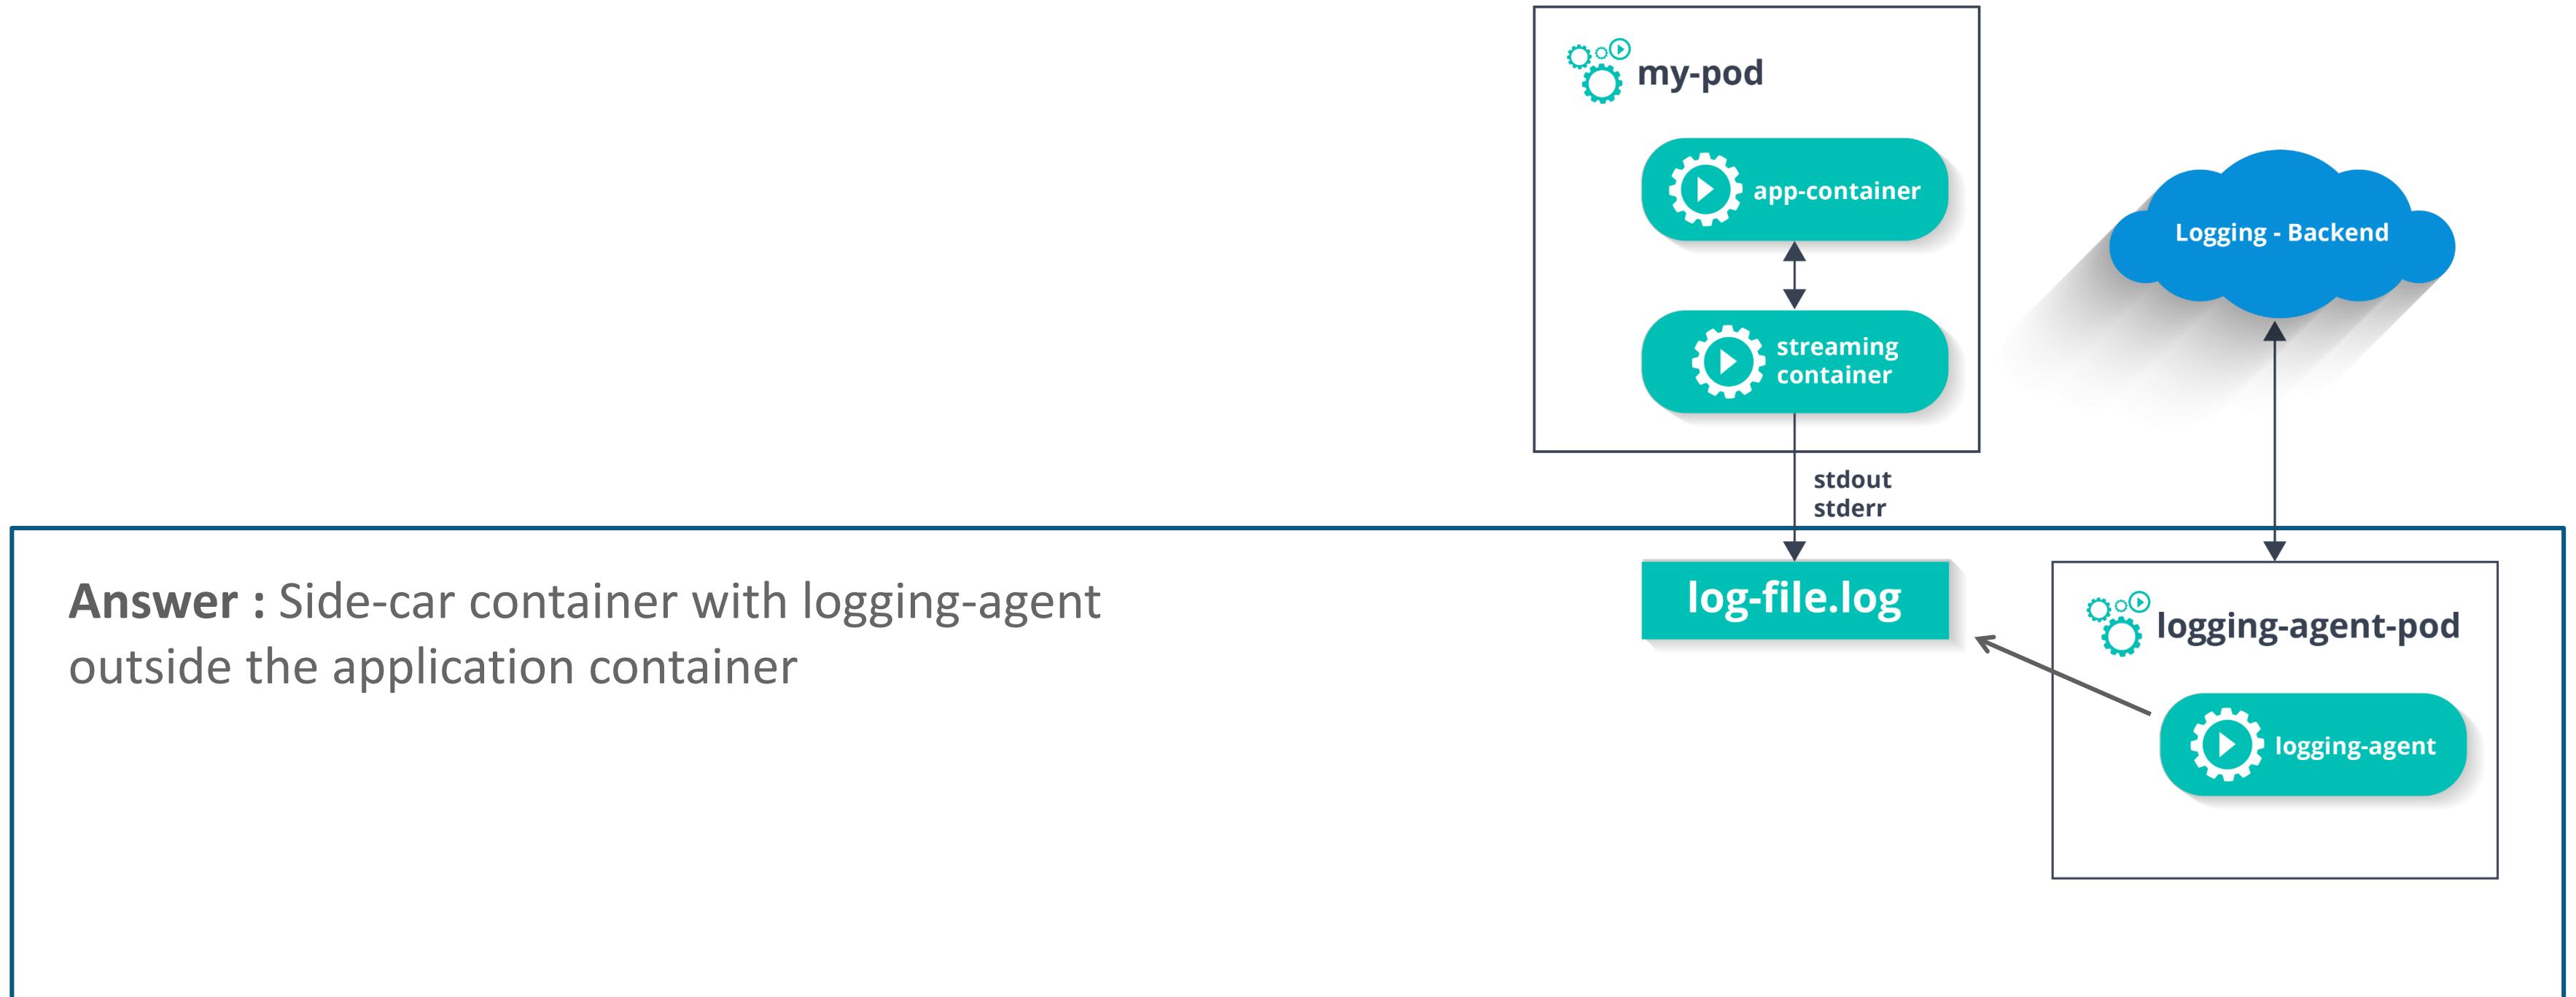
# Quiz

8. From the picture on side, explain the type of architecture it is ?



# Answers

8. From the picture on side, explain the type of architecture it is ?



# Summary

---

- In this module, you should have learnt:
  - Master Components of Kubernetes
  - Node components of Kubernetes
    - Docker, kubelet, kube-proxy, kubectl
  - Addons
    - Cluster DNS
    - Kubernetes Dashboard
    - Container Resource Monitoring
    - Cluster level logging



# Questions



# Thank You



For more information please visit our website  
[www.edureka.co](http://www.edureka.co)